

# Introduction to the AMD Versal™ AI Engine Architecture

2024.1

# Objectives

After completing this module, you will be able to:

## OBJECTIVE 01

Identify the major blocks in the AMD Versal™ architecture

## OBJECTIVE 02


Describe the architecture of the Versal AI Engine as well as its interfaces

## OBJECTIVE 03

Describe the AI Engine memory modules and memory access structure



# Versal Architecture: Overview



## Connectivity



- PCIe Gen5 & CCIX**
- High PCIe and DMA bandwidth
  - Cache coherency



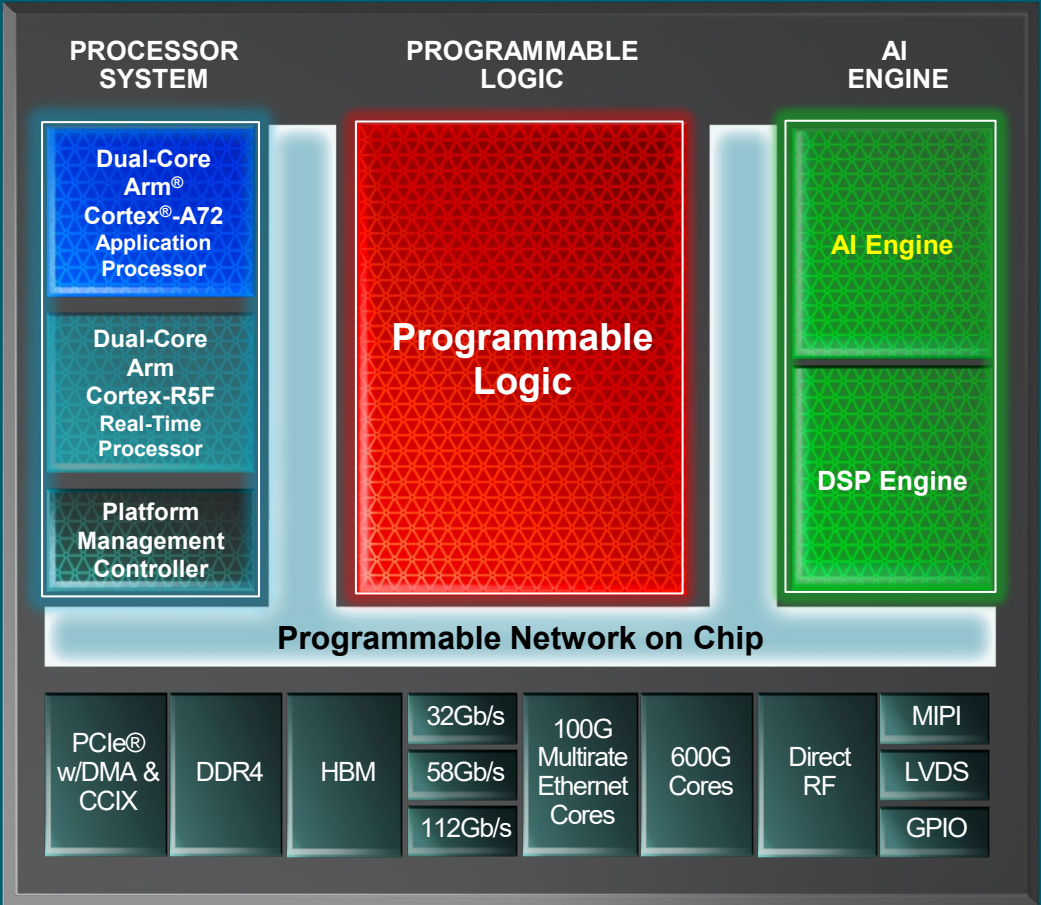
- Protocol Engines**
- 400G/600G cores
  - Power optimized



- Programmable I/O**
- Any interface or sensor
  - Includes 3.2Gb/s MIPI



- Transceiver Leadership**
- Broad range, 25G → 112G
  - 58G in mainstream devices




## Various Domain Connections



- Programmable NoC**
- Guaranteed bandwidth
  - Enables SW programmability

## Memory Options



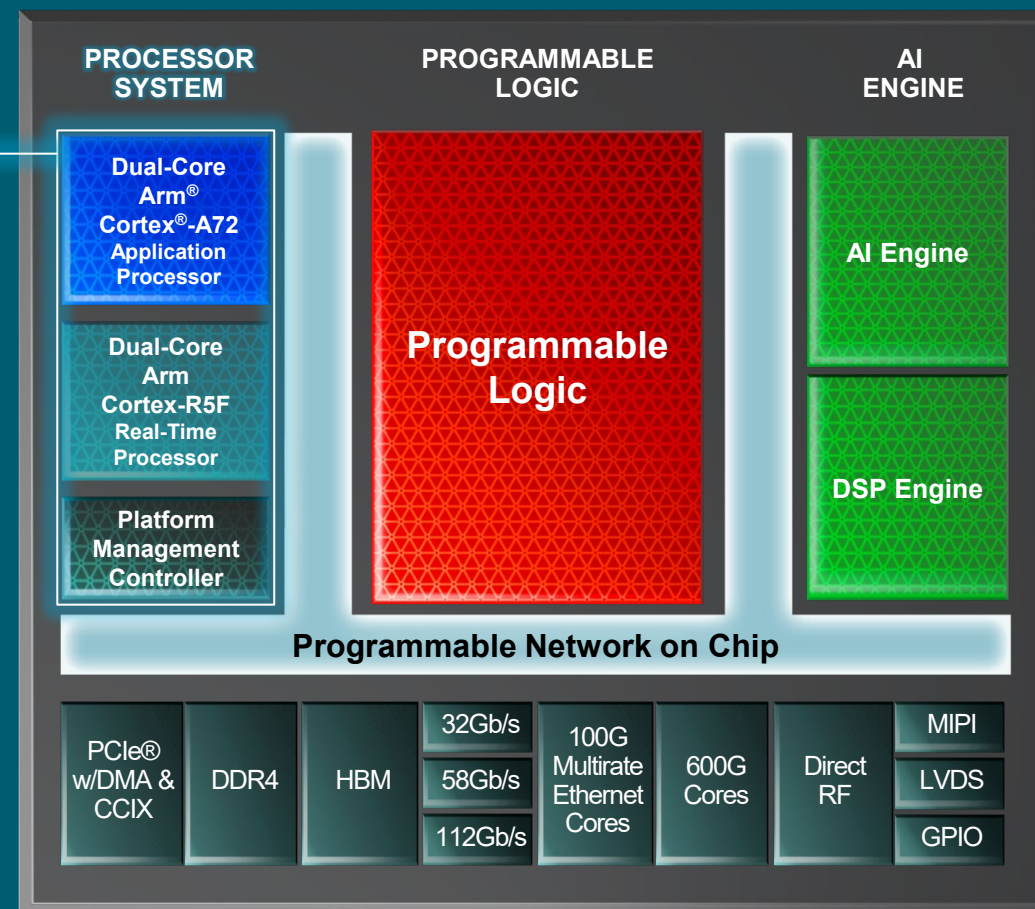
- DDR4 Memory**
- 3200-DDR4, 4266-LPDDR4
  - High bandwidth/pin

# Versal Architecture: Overview



## Processor System

- Execute complex algorithms and **decision making**
- Provide safety processing and redundancy for mission- and safety-critical applications
- **Manage the entire platform**
- **Load** each aspect of the **Versal™** device and **monitor status**
- Support capability extension
  - PL-instantiated MicroBlaze™ processor

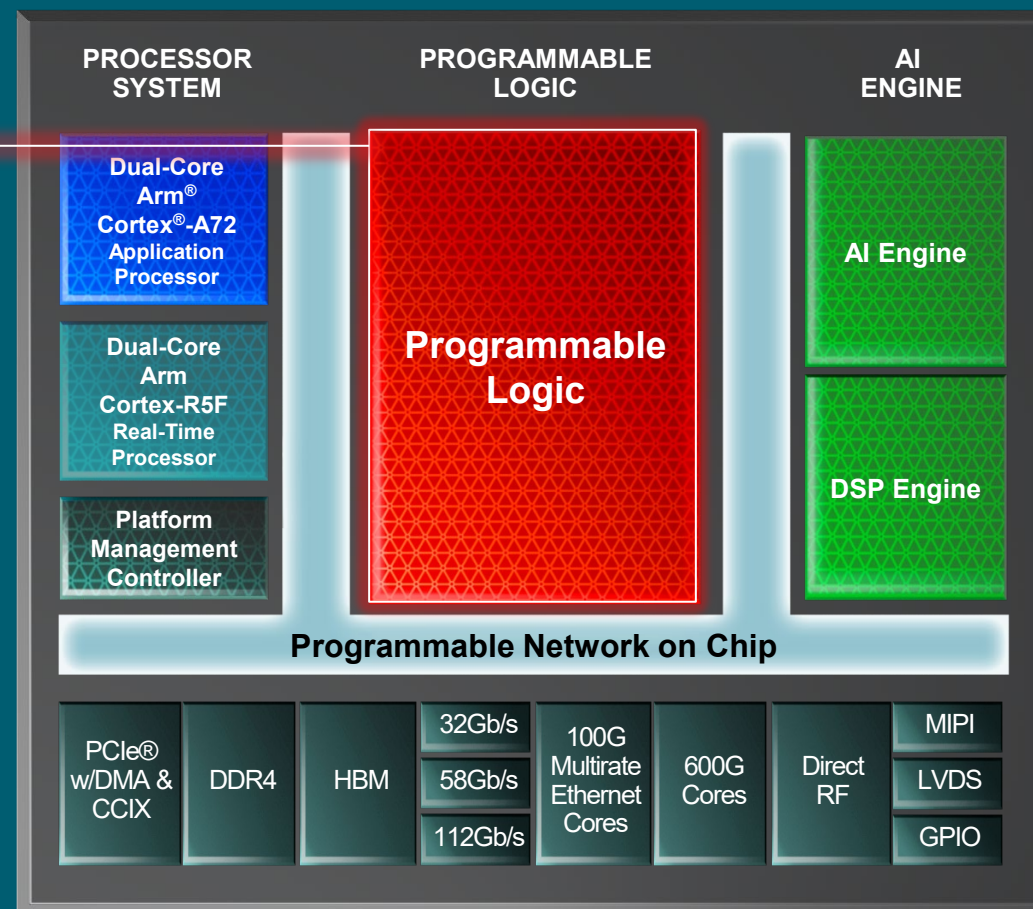


# Versal Architecture: Overview



## Programmable Logic

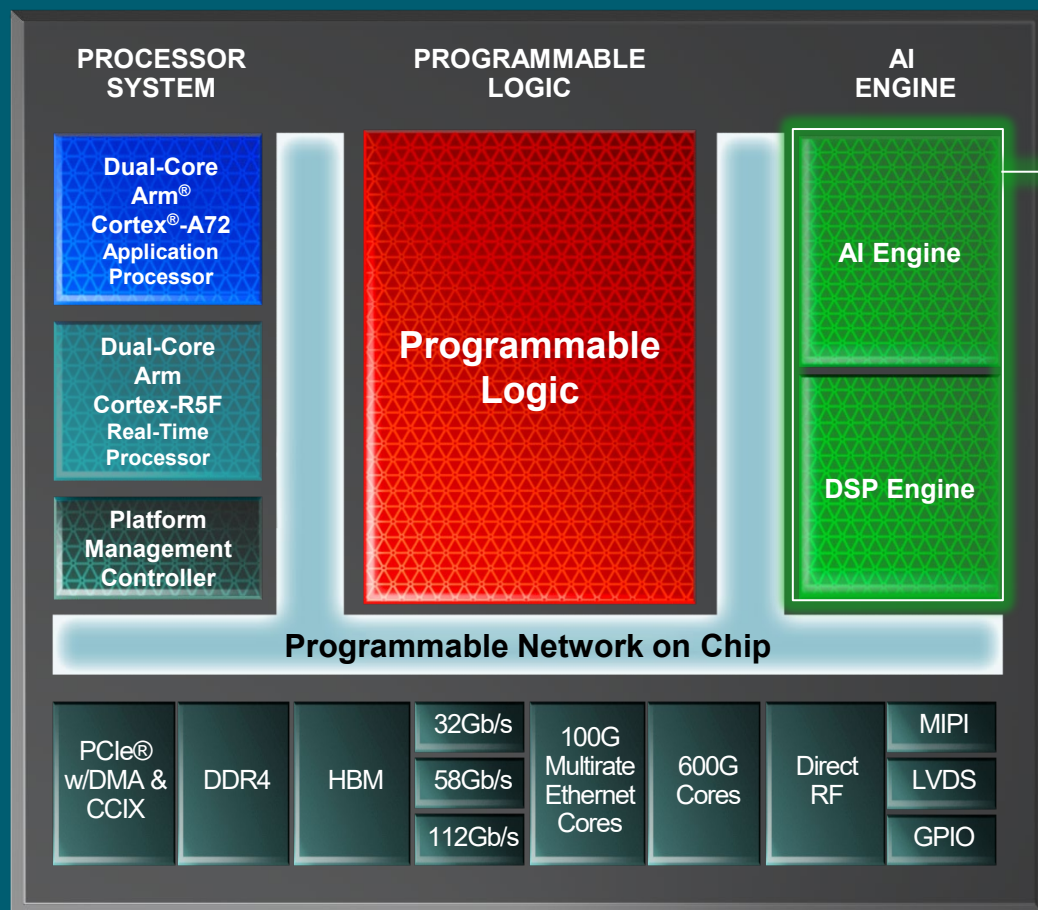
- Millions of reconfigurable system logic cells
- Support for parallel, pipelined, and hybrid architectures
- Wide variety of memory elements





# Versal Architecture: Overview

*This is not a physical representation!*



## AI Engine

- AI compute
- Diverse DSP workloads

## AI Is Everywhere

- Wired communications, automotive, and consumer markets

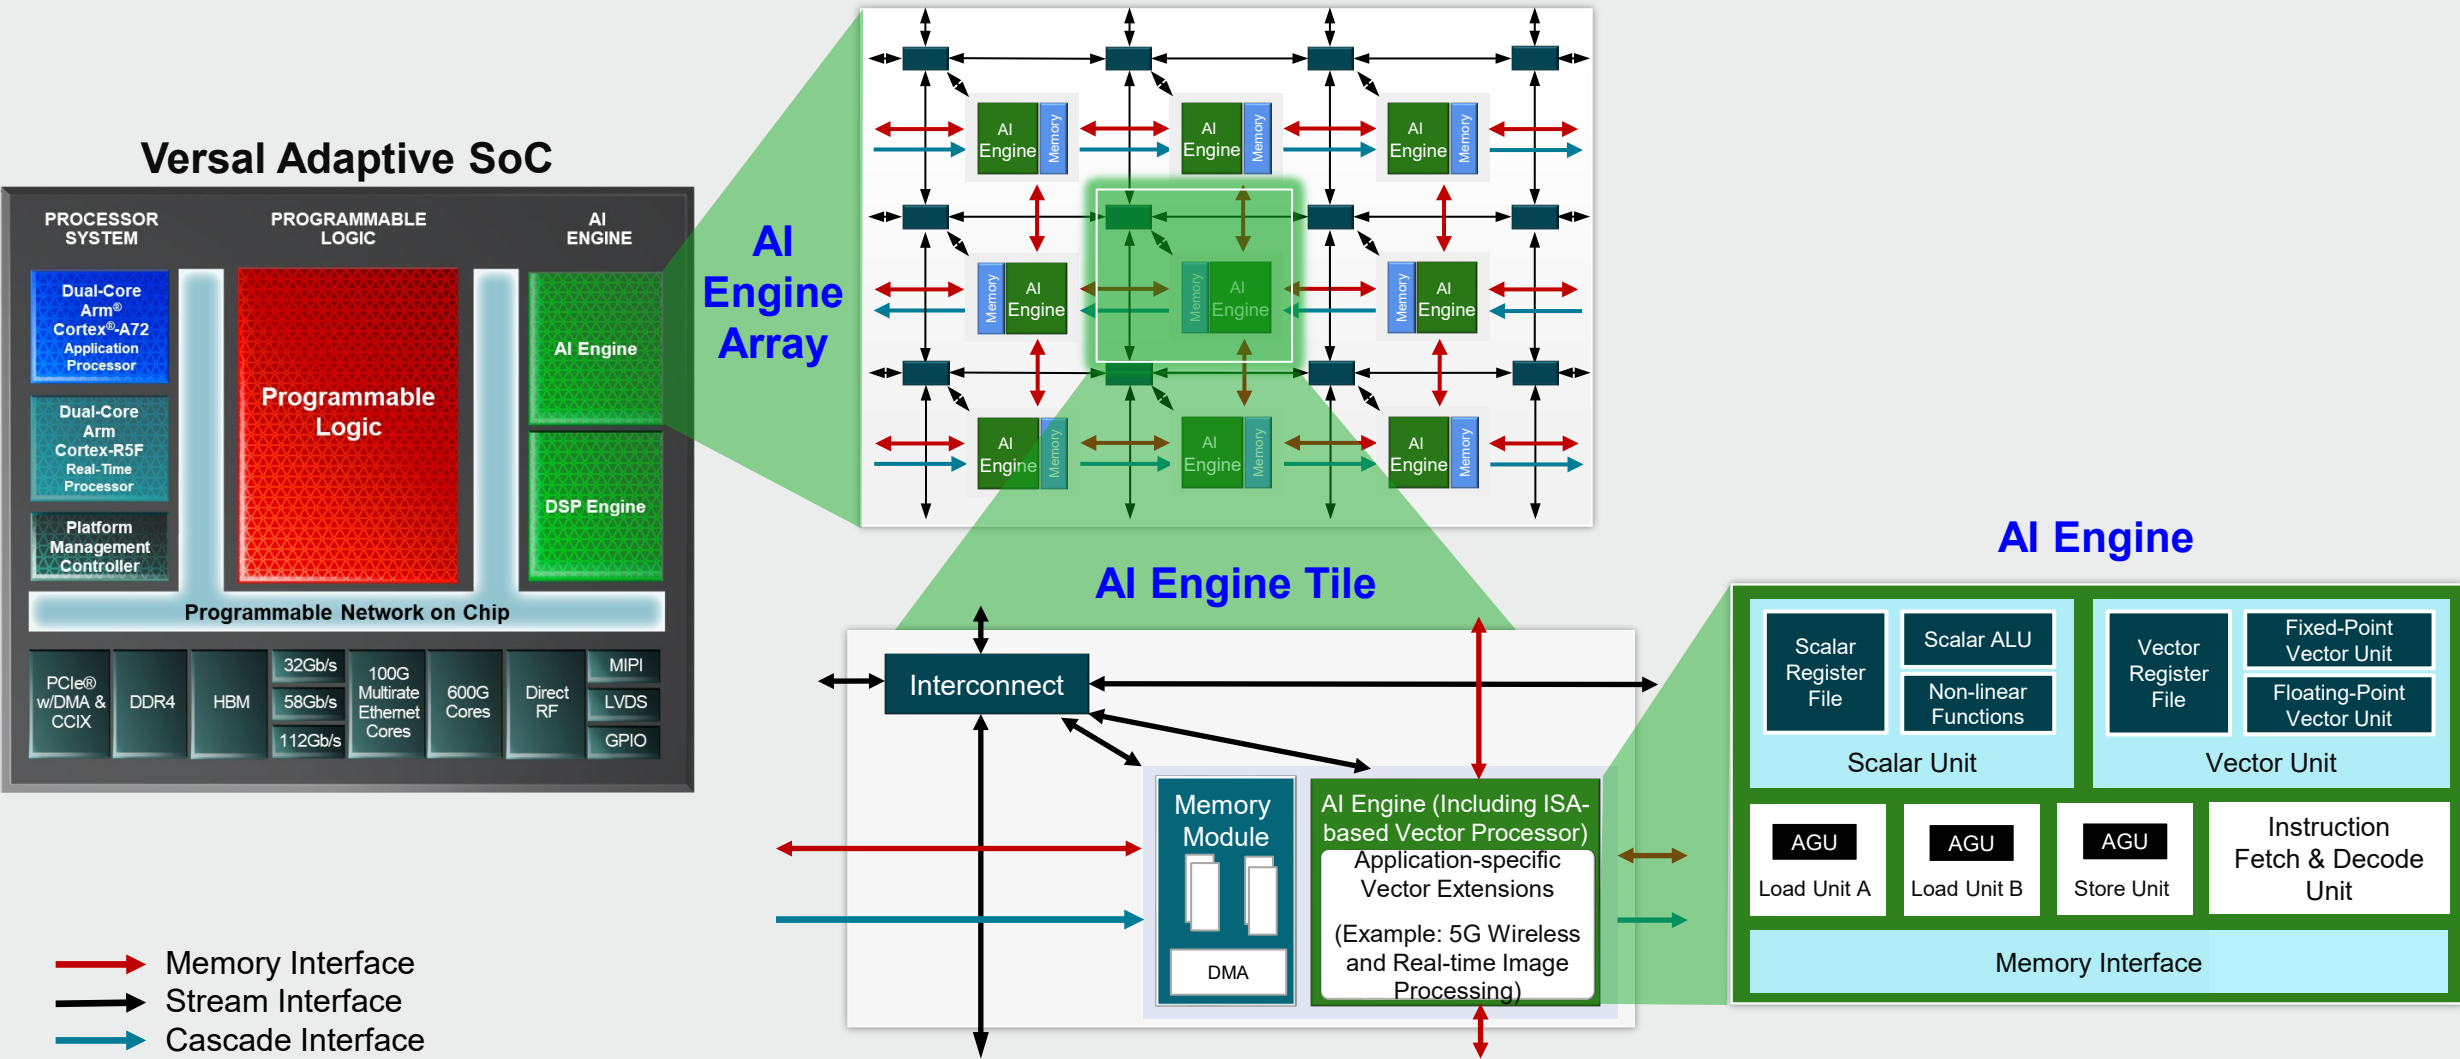
## DSP Engines

- High-precision, floating-point computation support
- Offload additional functions for acceleration

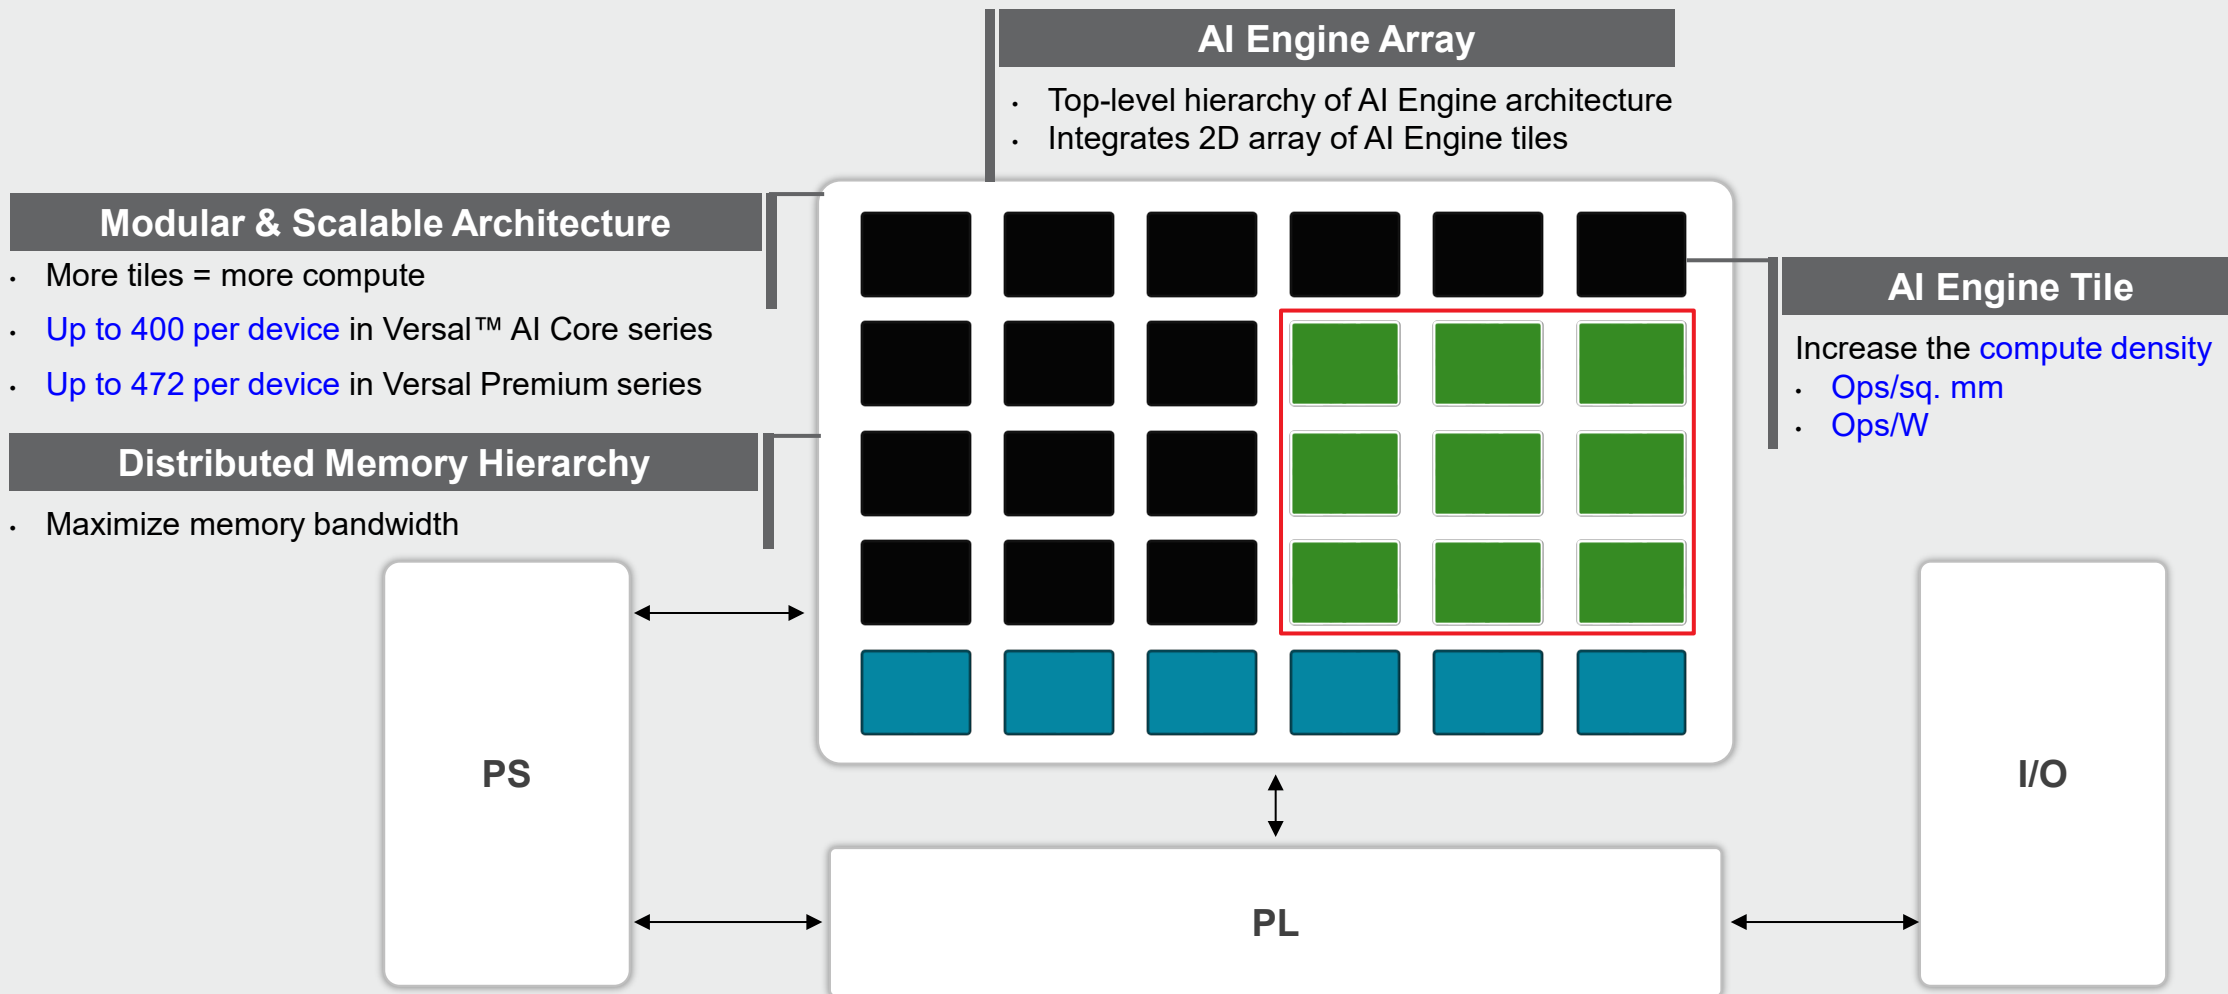
## AI Engines

- High throughput, low latency, and power efficient
- Ideal for AI inference and advanced signal processing
- Array of VLIW processors with SIMD vector units
- Instruction-level and data-level parallelism

# AI Engine: Terminology




# AI Engine: Array Architecture

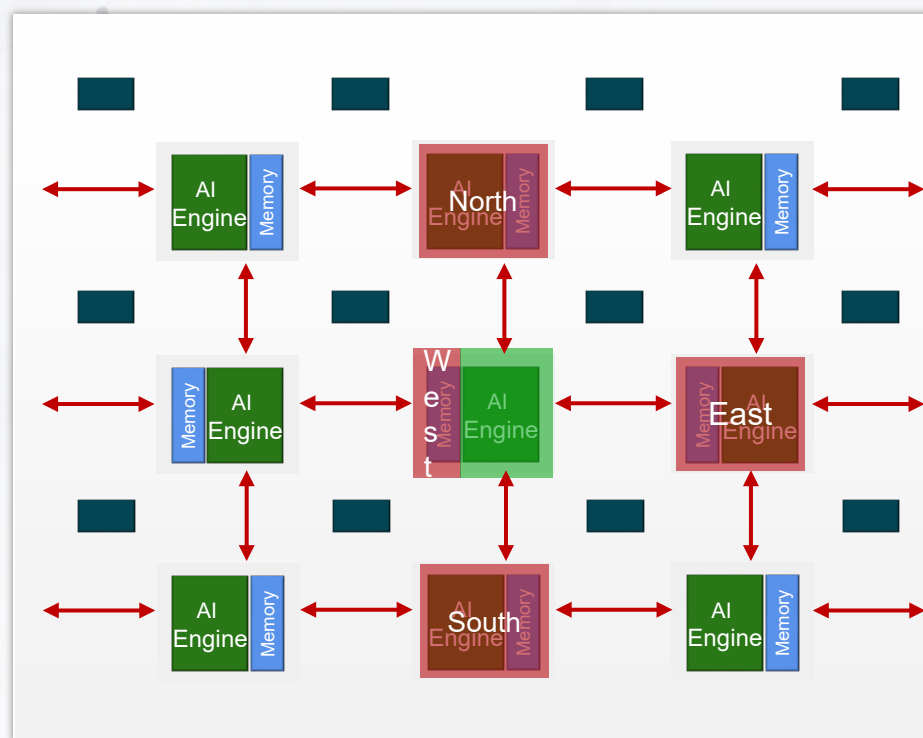




# AI Engine: Array Architecture




3 x 3 array of AI Engines tiles



Memory Interface

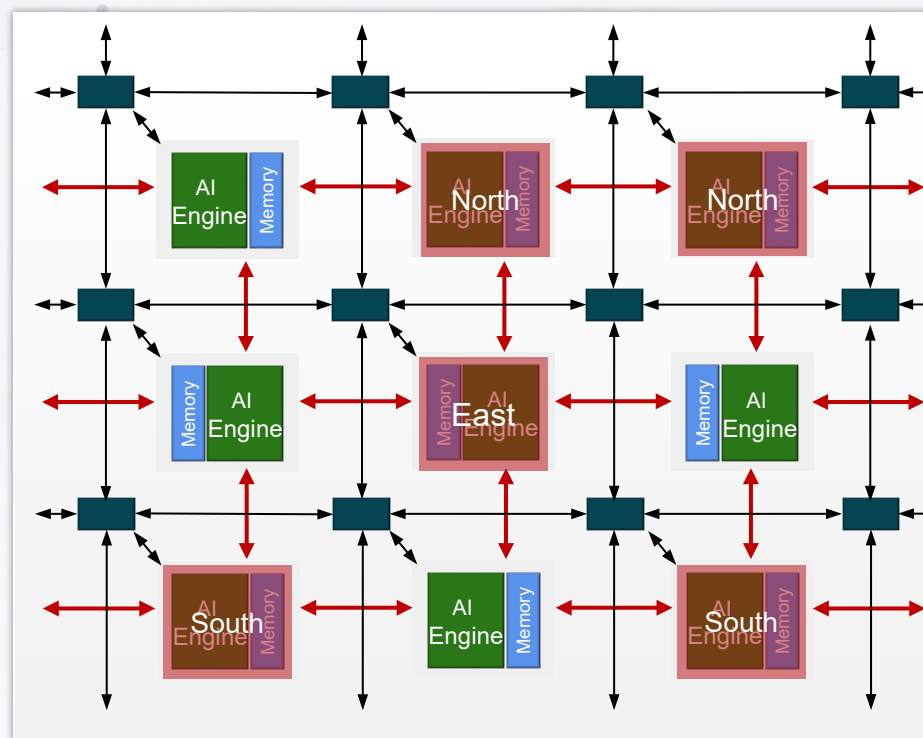
The tile can communicate to the local memory on each neighboring tile: north, south, east, or west

# AI Engine: Array Architecture




The non-blocking interconnect allow data communications between non-neighboring tiles throughout the array.

3 x 3 array of AI Engines tiles



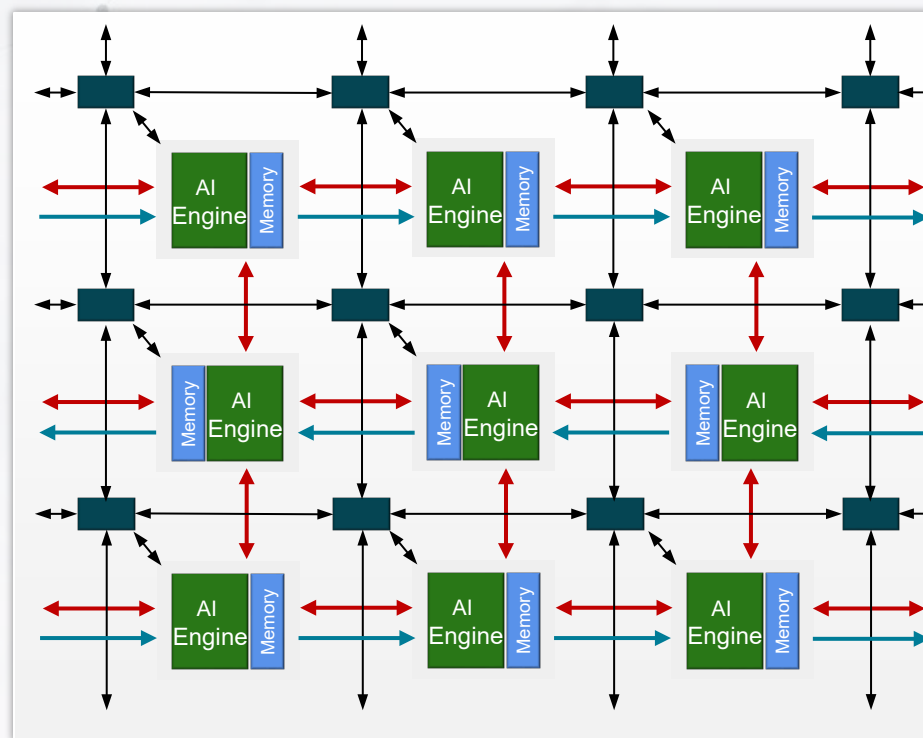
Memory Interface  
Stream Interface

# AI Engine: Array Architecture



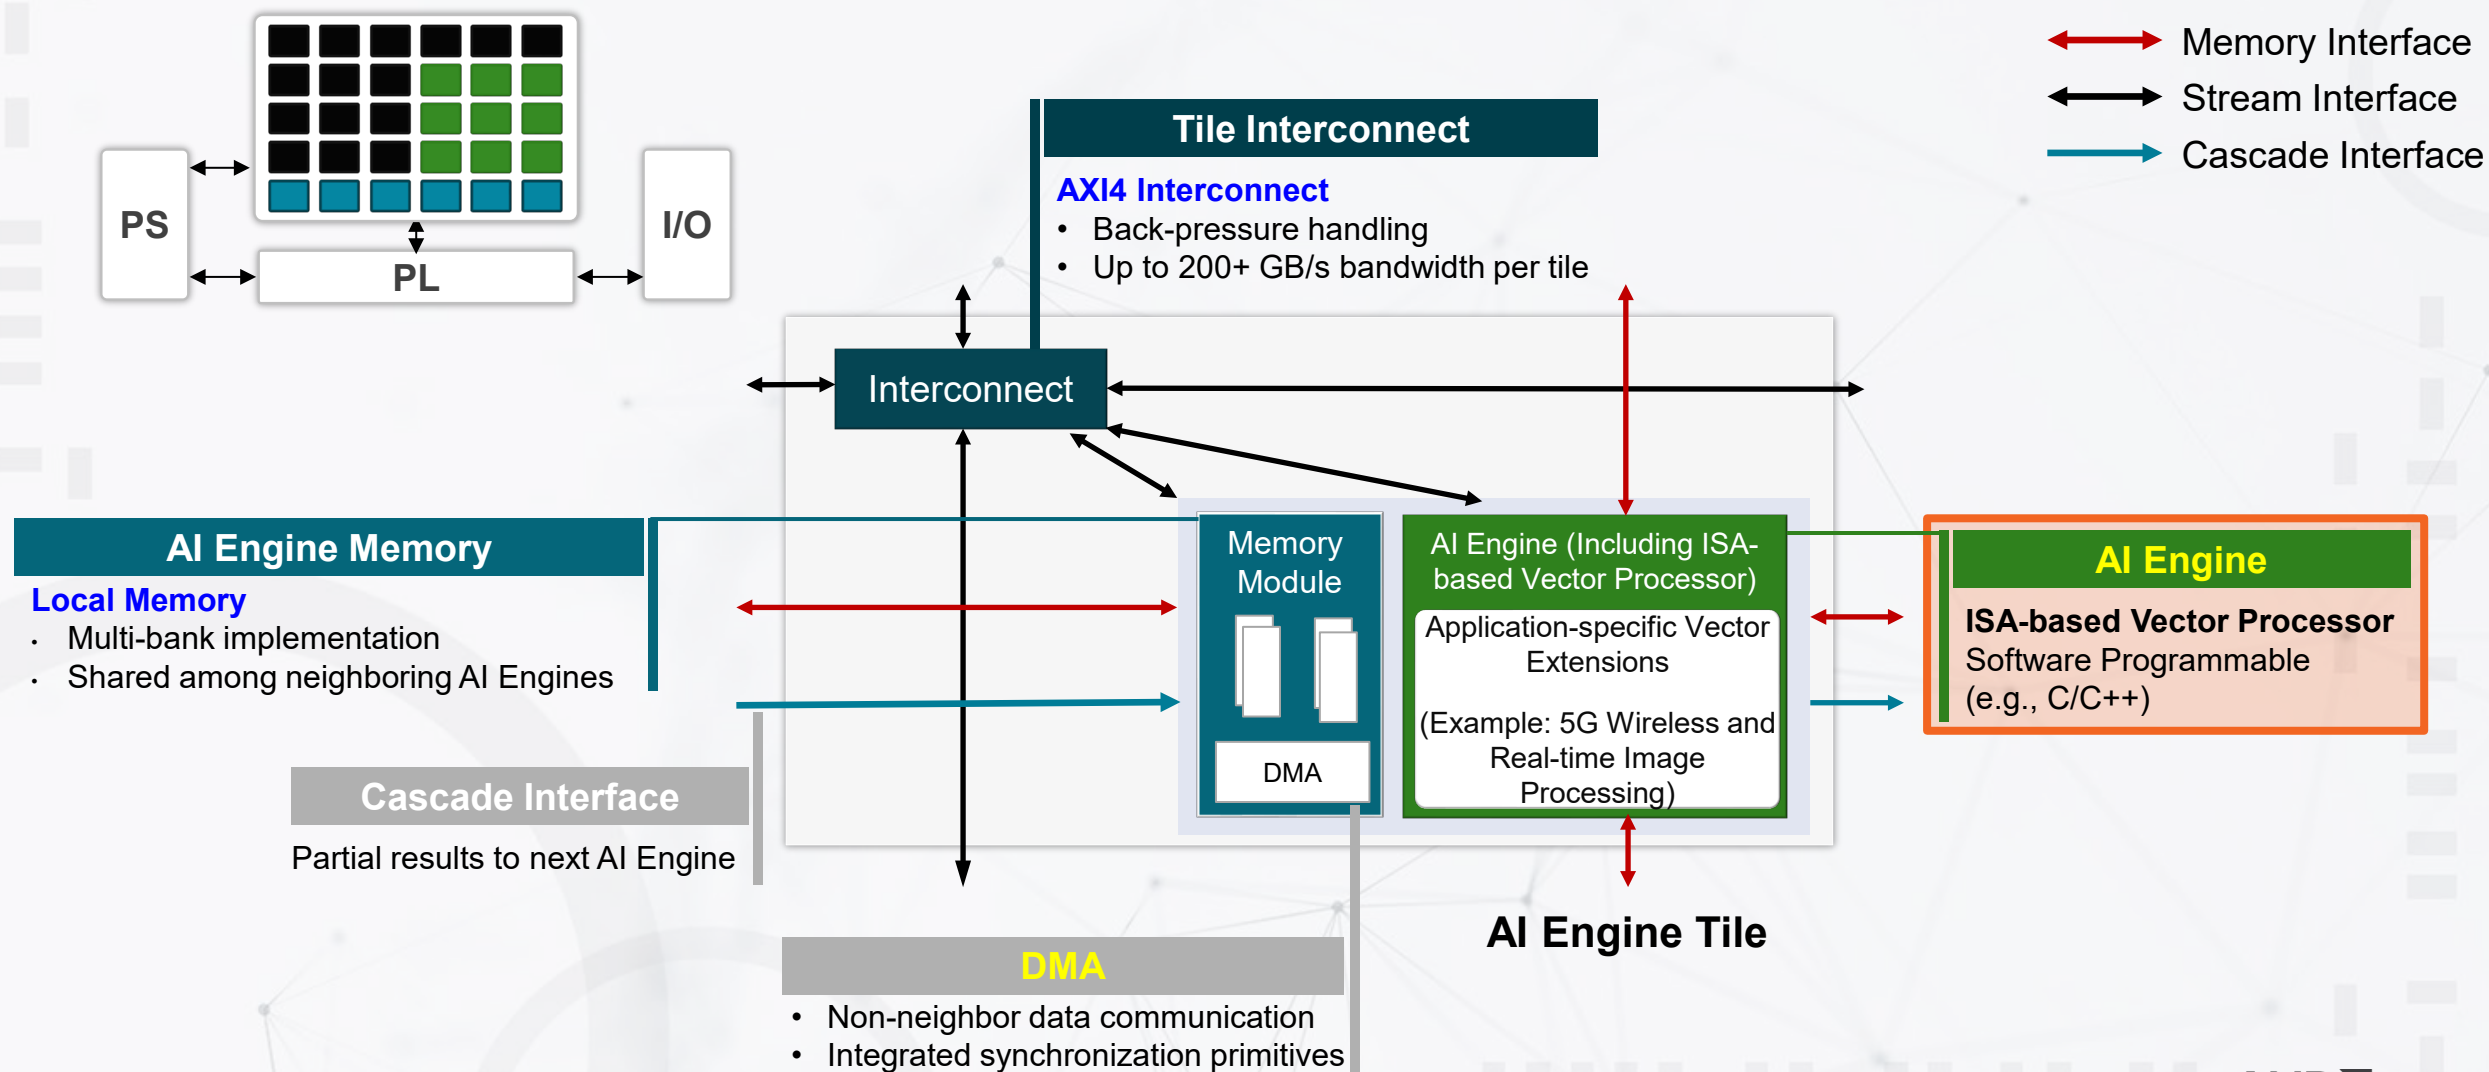
Similarly, the blue arrows indicate the cascading of partial compute to neighboring tiles.

3 x 3 array of AI Engines tiles

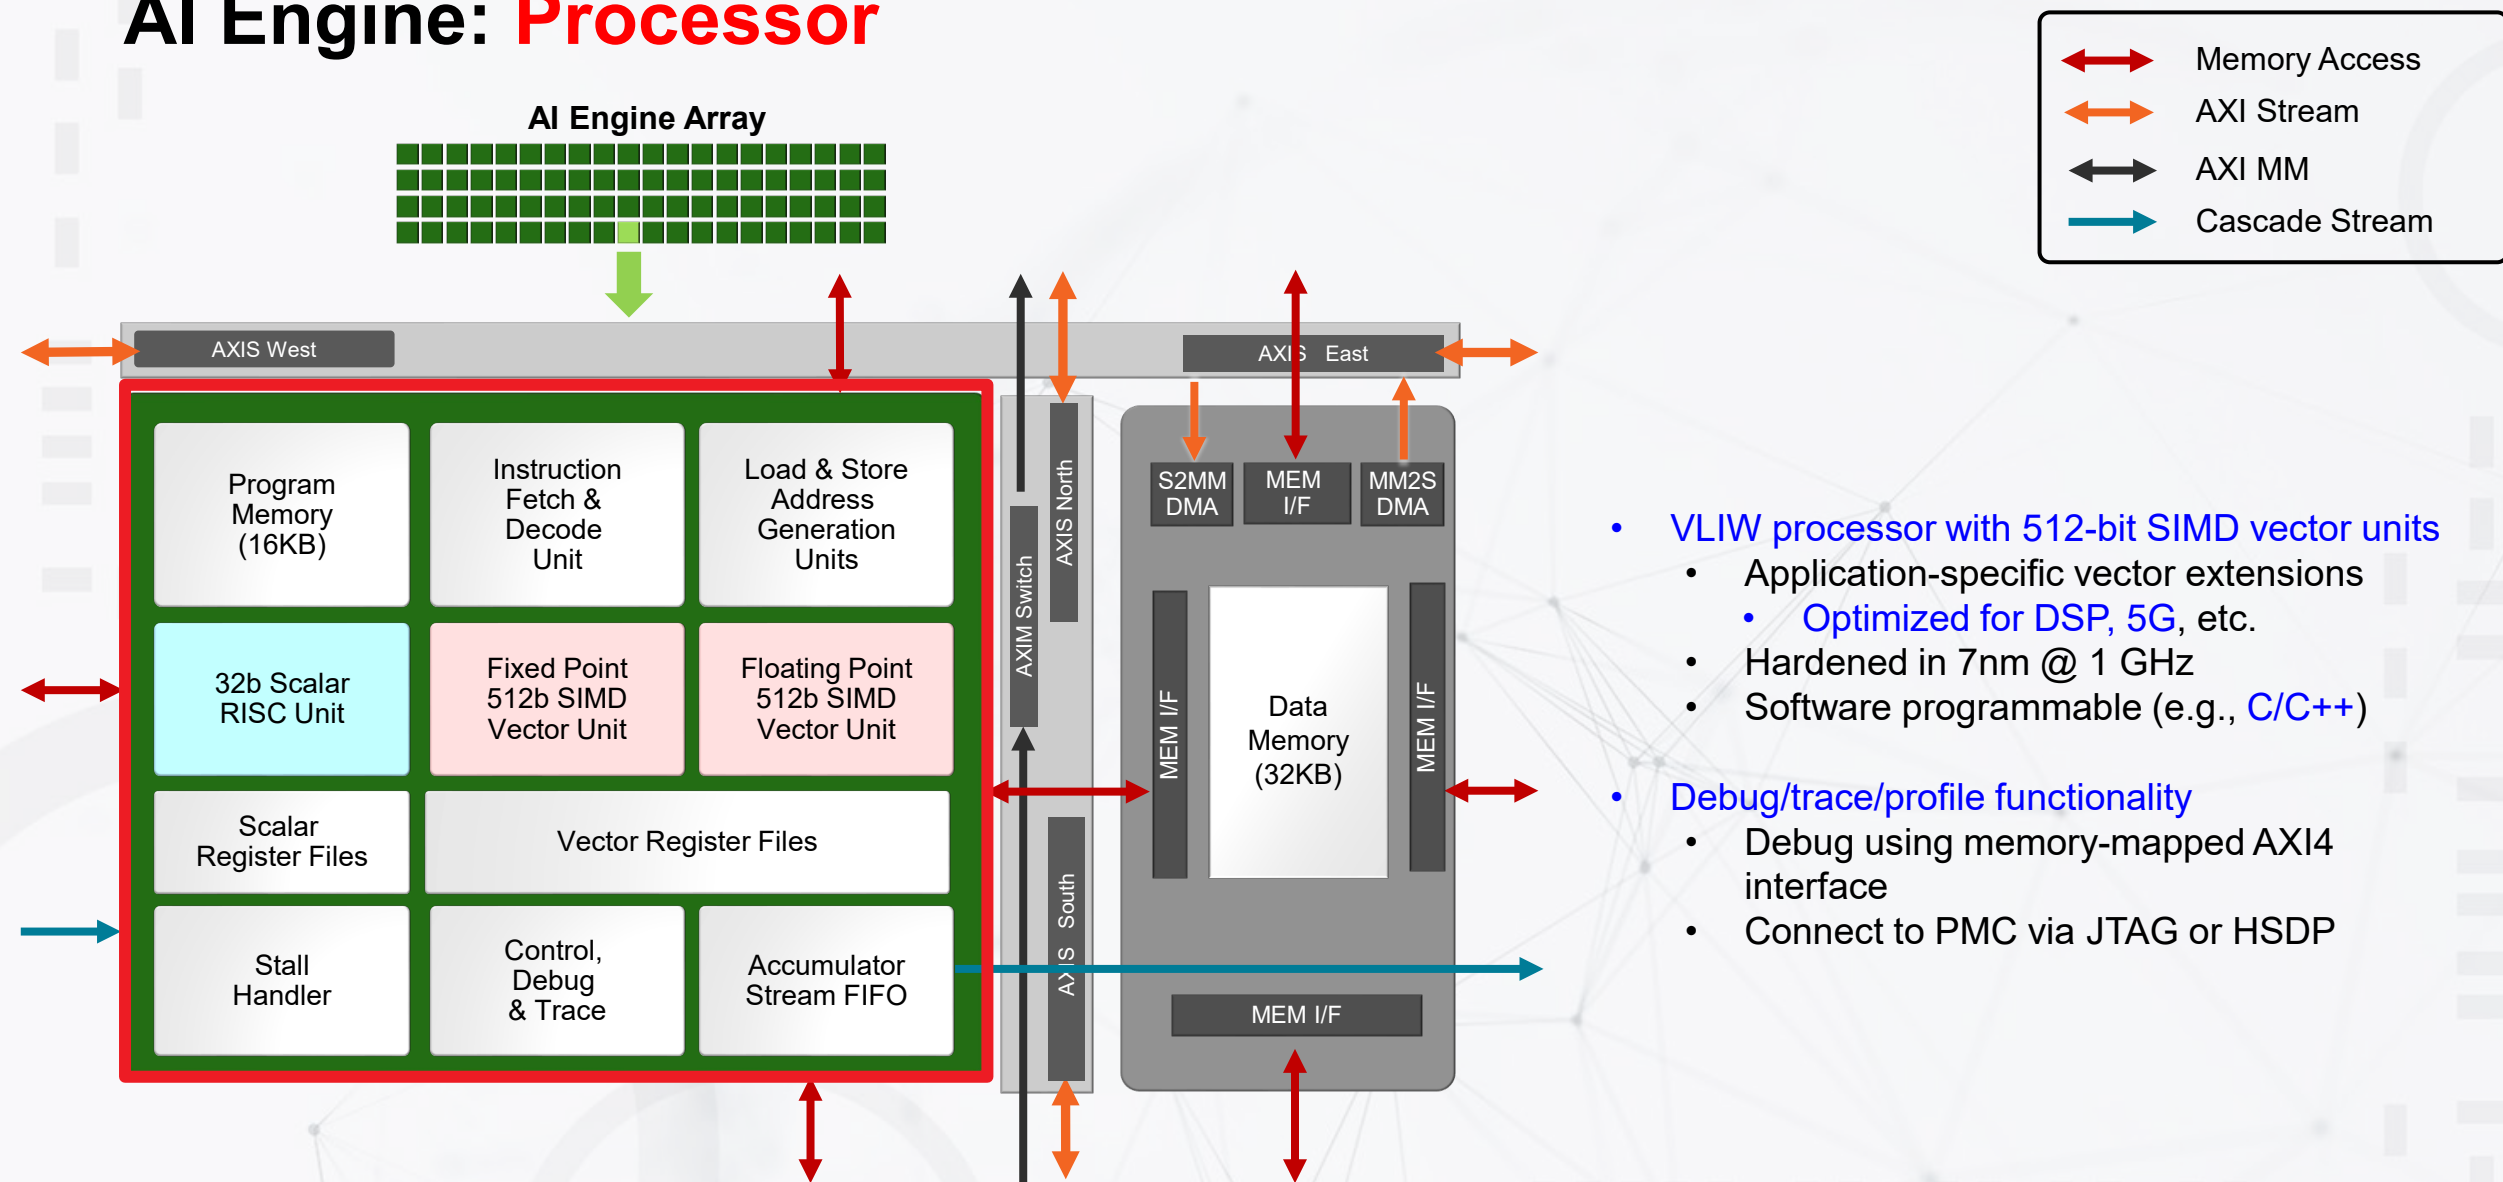


Memory Interface  
Stream Interface  
Cascade Interface

# AI Engine: **Tile-based Architecture**



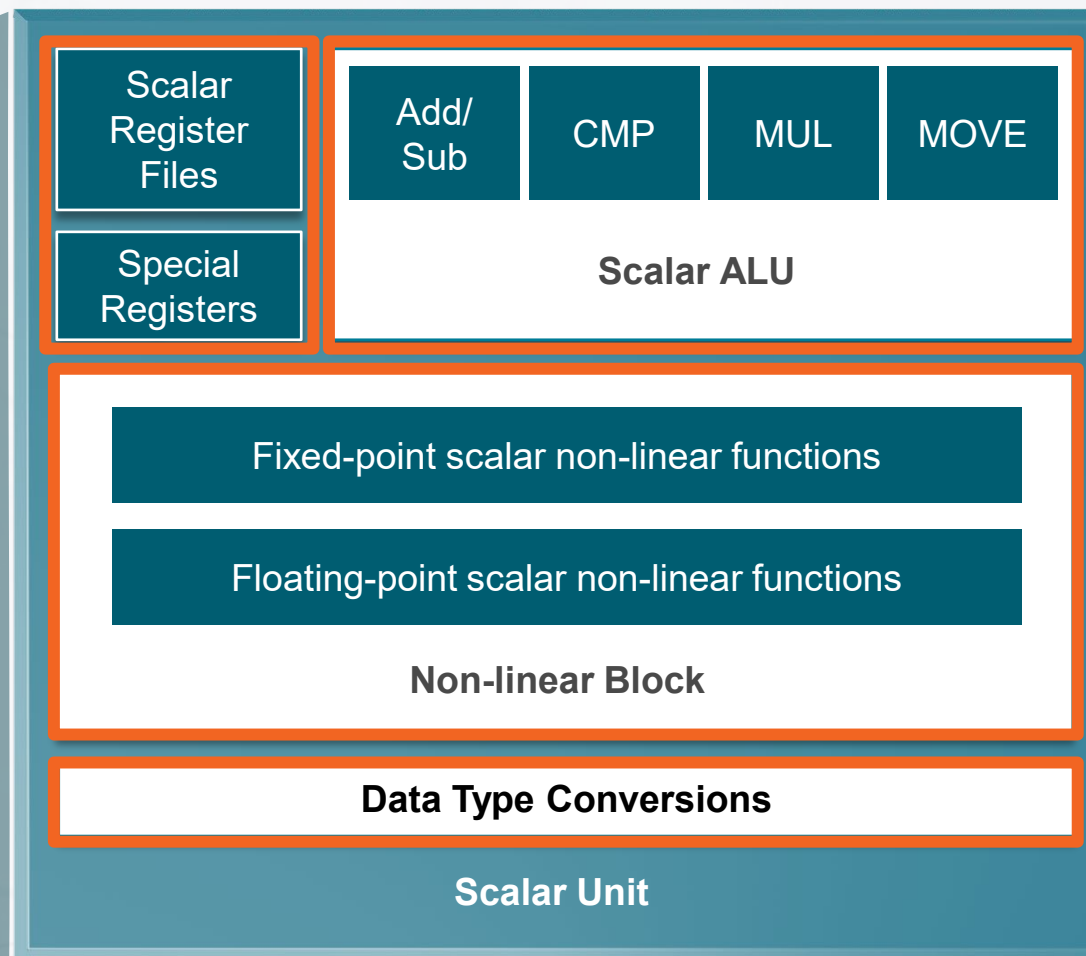
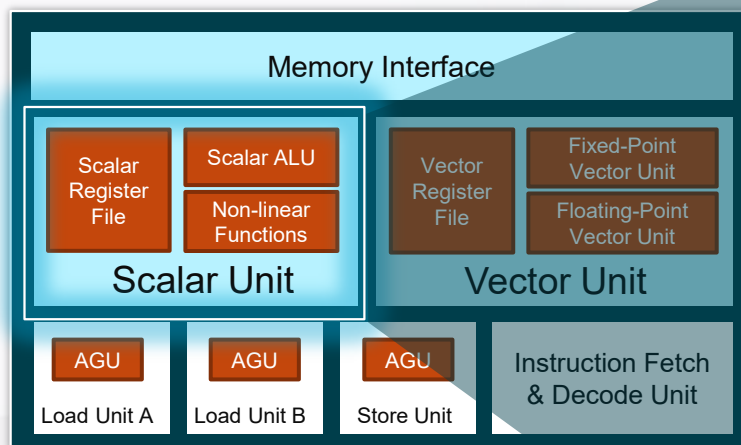
# AI Engine: Processor



- VLIW processor with 512-bit SIMD vector units
  - Application-specific vector extensions
    - Optimized for DSP, 5G, etc.
  - Hardened in 7nm @ 1 GHz
  - Software programmable (e.g., C/C++)
- Debug/trace/profile functionality
  - Debug using memory-mapped AXI4 interface
  - Connect to PMC via JTAG or HSDP



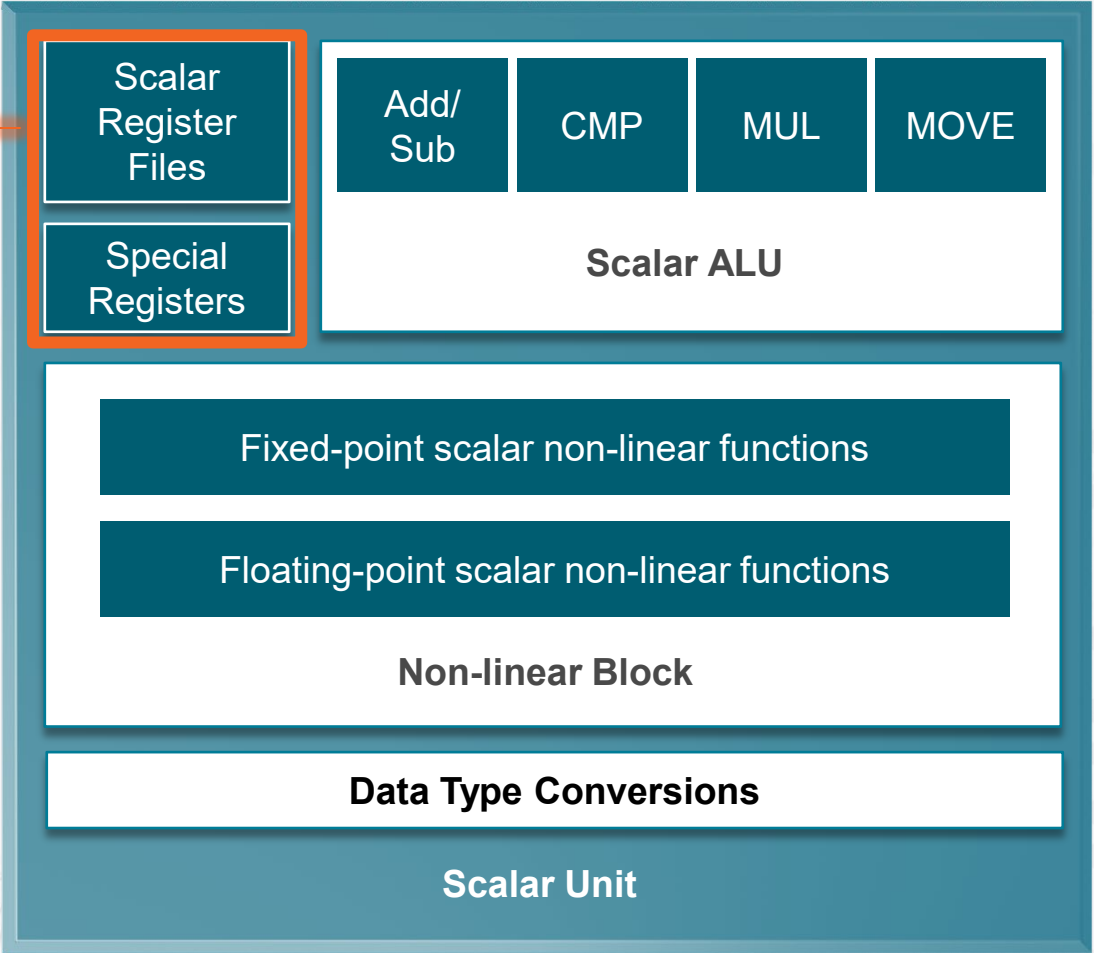
# AI Engine – Scalar Unit



# Scalar Unit – Registers

Scalar registers include configuration registers

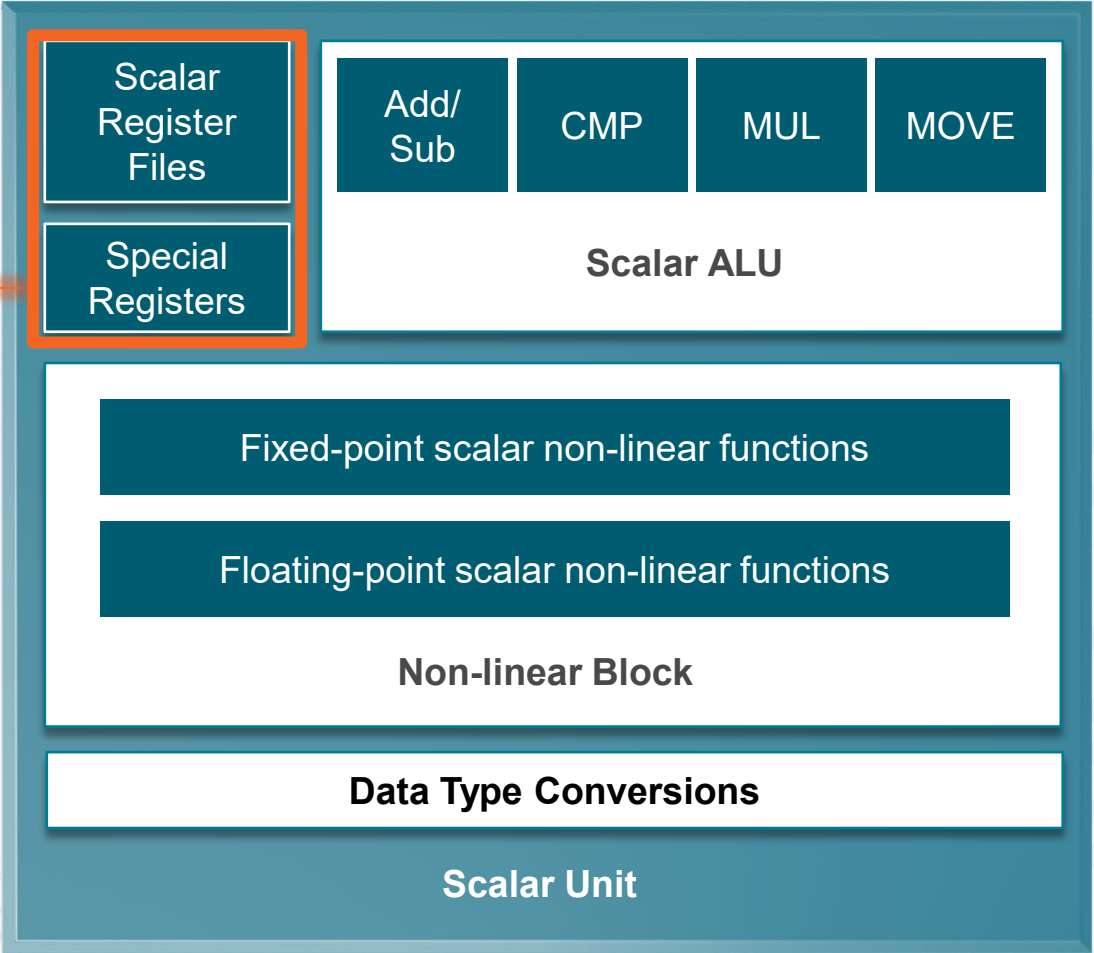
Syntax	No. of Bits	Description
r0...r15	32 bits	General-purpose registers
m0...m7	20 bits	Modifier registers
p0...p7	20 bits	Pointer registers
cl0...cl7	32 bits	Configuration registers
ch0...ch7		
c0...c7	64 bits	



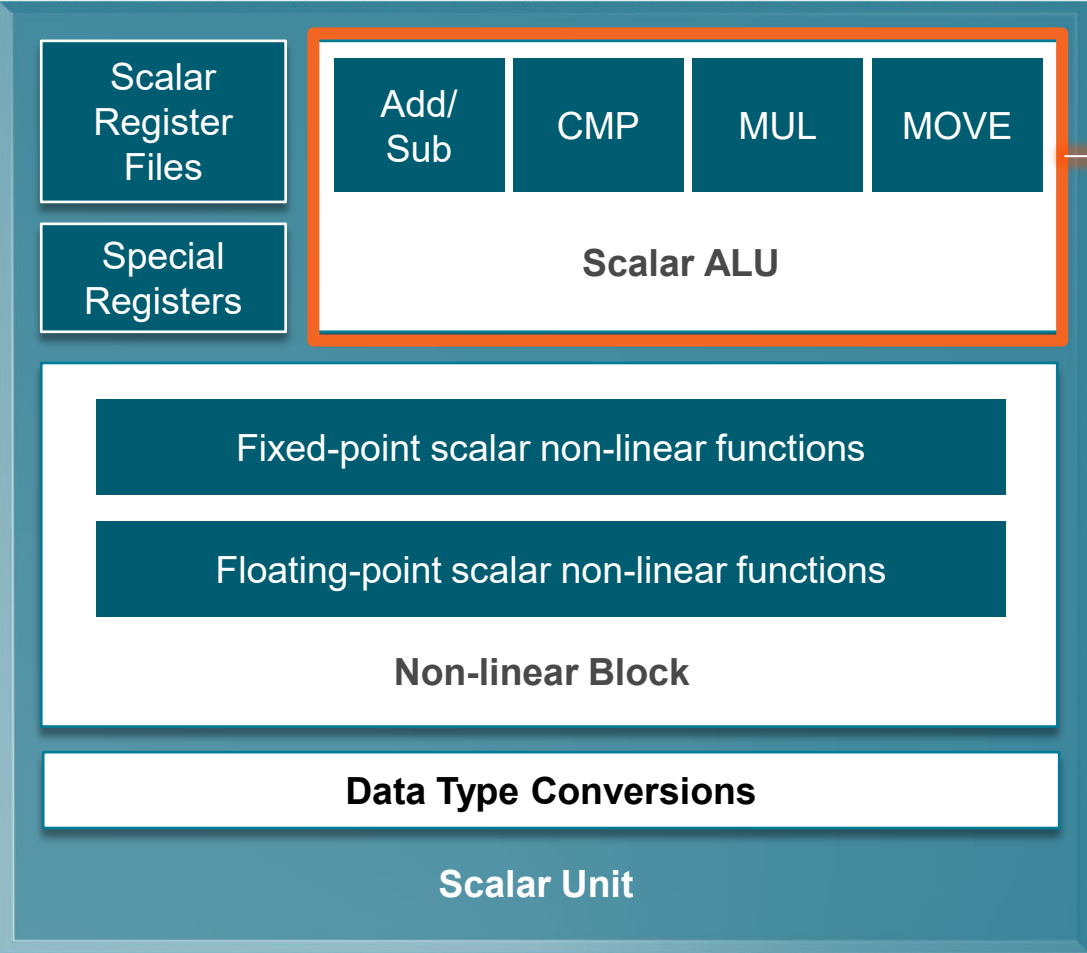
# Scalar Unit – Registers

## Special registers

Syntax	No. of Bits	Description
cb0...cb7	20 bits	Circular buffer start address
cs0...cs7	20 bits	Circular buffer size
wcs0...wcs3	40 bits	Wide circular buffer size
s0...s7	8 bits	Shift control
sp	20 bits	Stack pointer
lr	20 bits	Link register
pc	20 bits	Program counter
fc	20 bits	Fetch counter
mc0...mc1	32 bits	Status register
md0...md1	32 bits	Mode control register
ls	20 bits	Loop start
le	20 bits	Loop end
lc	32 bits	Loop count
lci	32 bits	Loop count (PCU)
S	8 bits	Shift control



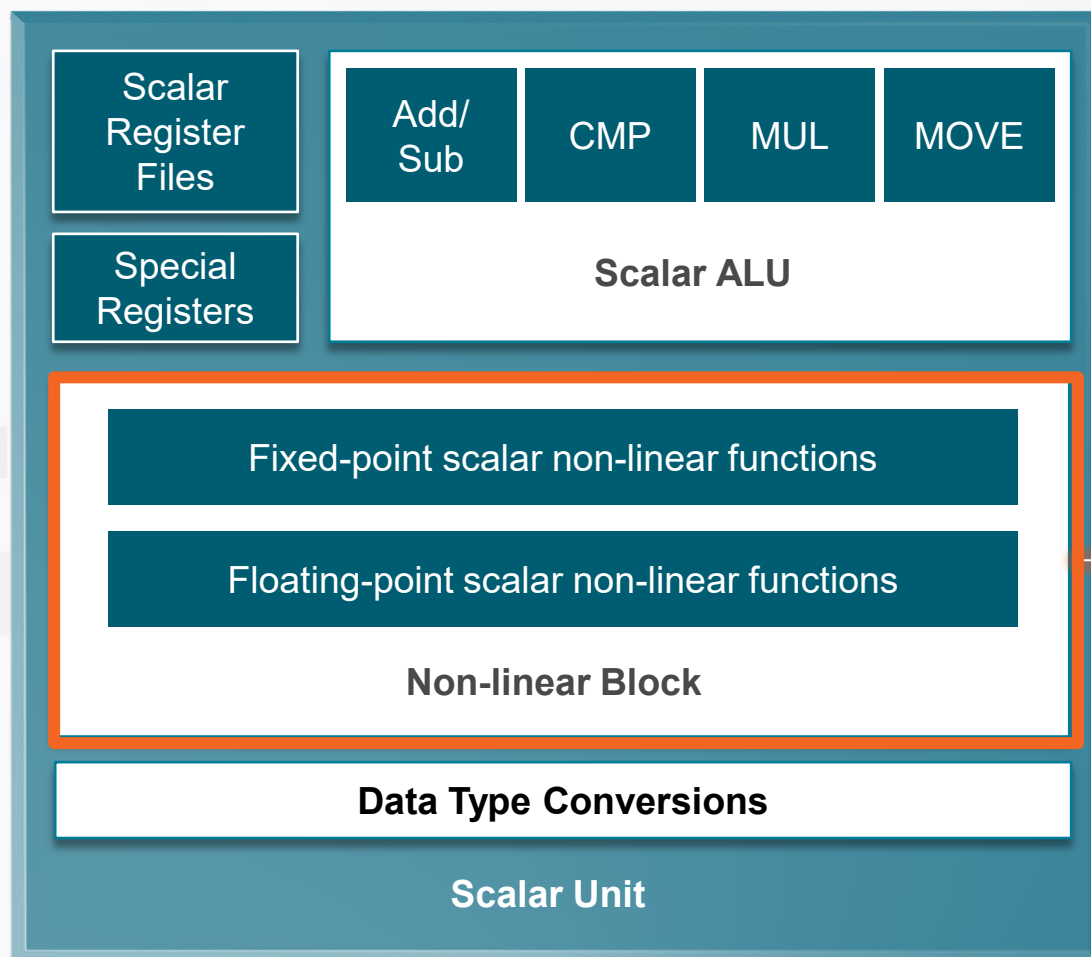
# Scalar Unit – ALU



AIE compiler supports standard C data types

Operations	Cycles
32-bit add/subtract	1
32-bit logical operation	1
32x32-bit mult (signed & unsigned)	3
Shift and compare	1

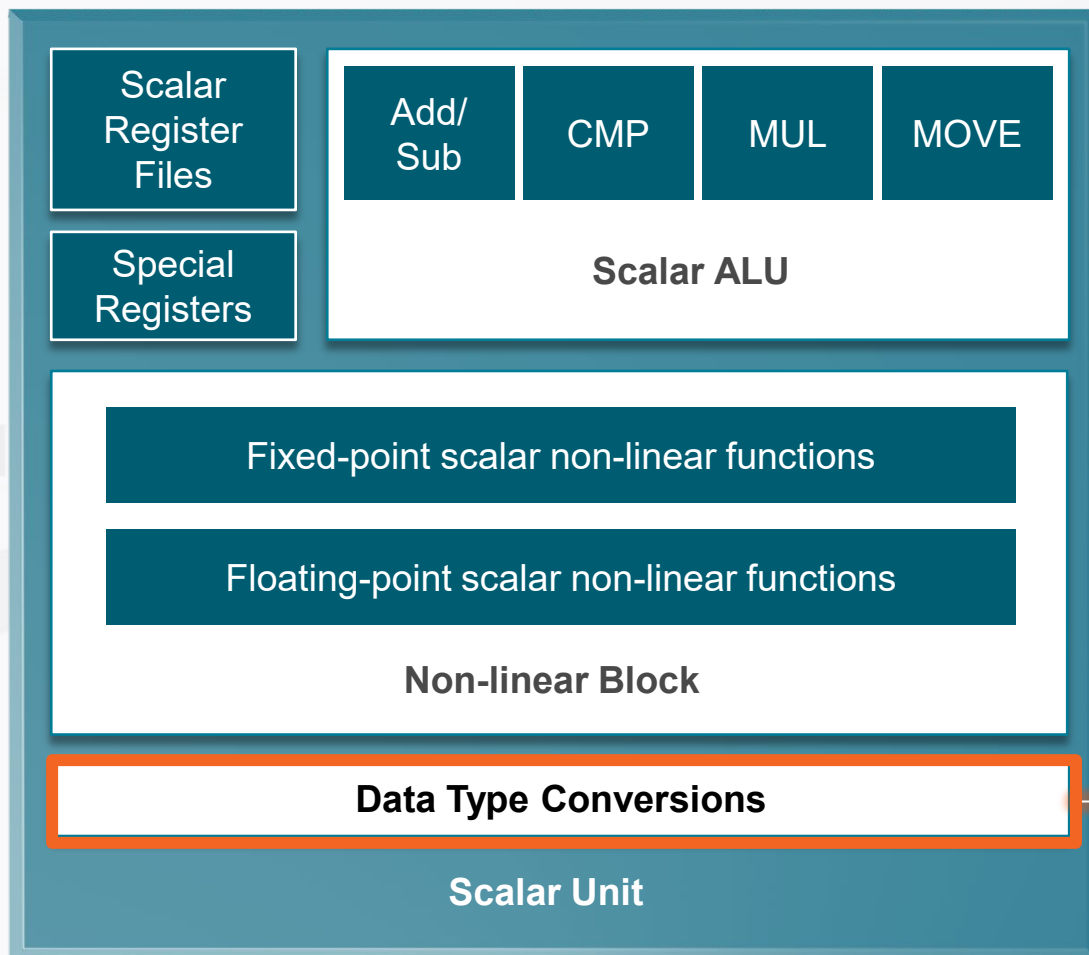
# Scalar Unit – Non-linear Blocks



- **Fixed point/integer non-linear functions**
  - sin/cos, abs (real input), clz, min/max
  - sqrt, inv sqrt & inv – supported via data conversion
    - Fix2float and float2fix
- **Floating-point, non-linear functions**
  - sqrt, inv sqrt, inv, abs, min/max, comparison



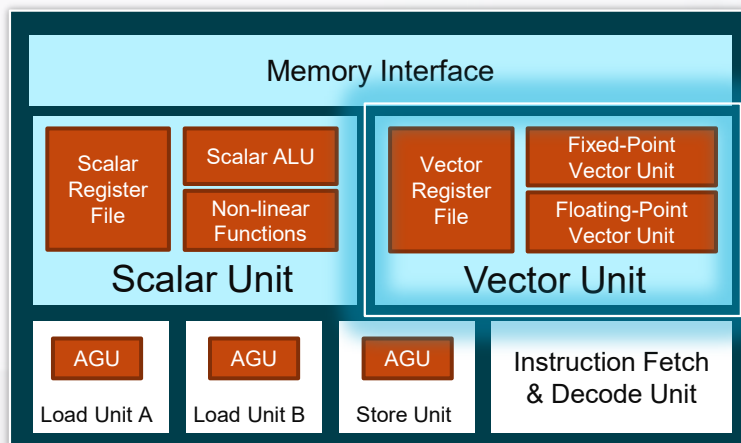
# Scalar Unit – Data Type Conversions



AIE compiler supports standard C data types

- short
- char
- int
- long
- float
- double

# Vector Unit



**1** Fixed-point Vector Unit

**2** Floating-point Vector Unit

**3** Vector Register File

**4** Multi-precision Support

# Fixed-point Vector Unit

## Multiply Accumulator (MAC) path

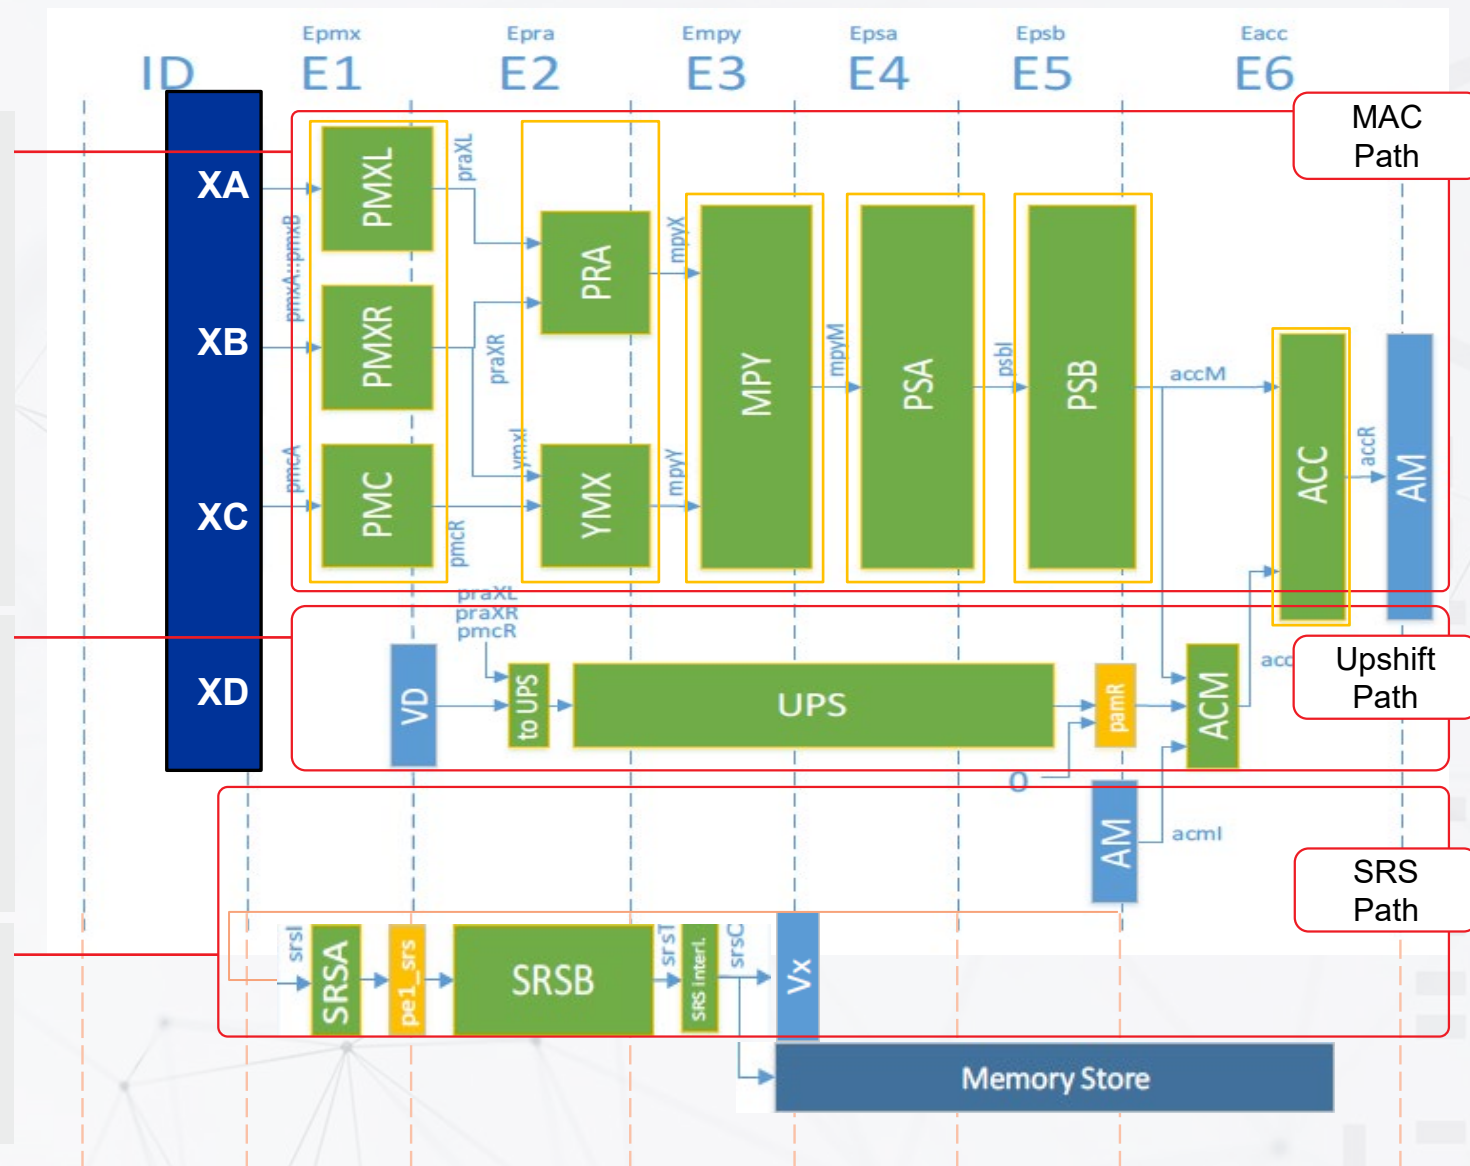
- Reading data from vector registers
- Full permute unit (32-bit granularity)
- Pre-adder
- Multiplier
- Two-step post-adding
- Accumulator

## Upshift path

- Reading from permute units/vector registers
- Left shifting
- Accumulator

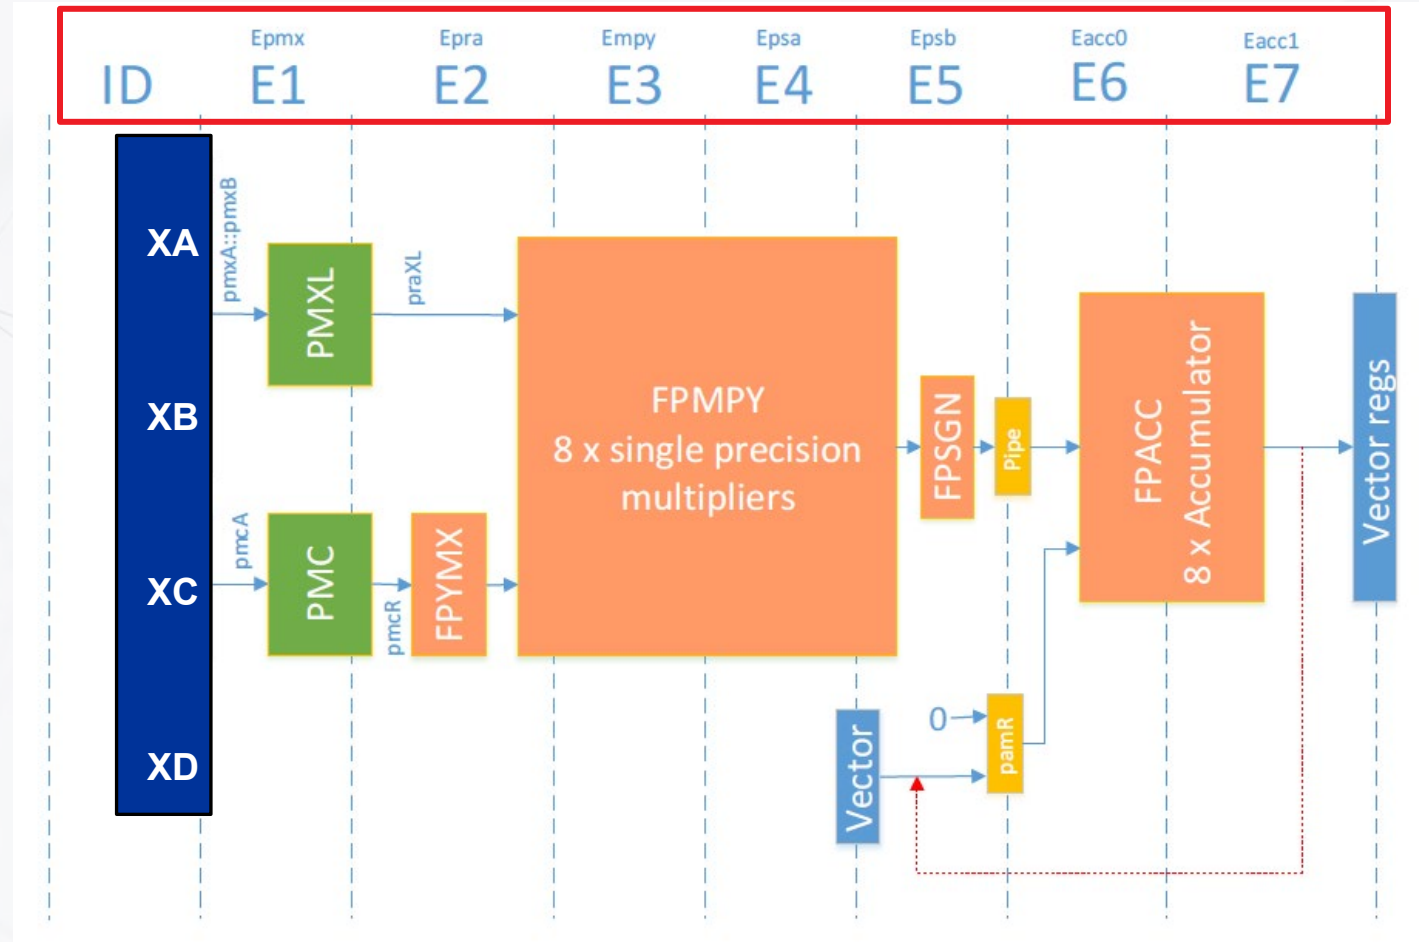
## Shift-Round-Saturate (SRS) path

- Reading from accumulator registers
- Storing to registers/memory



# Floating-point Vector Unit

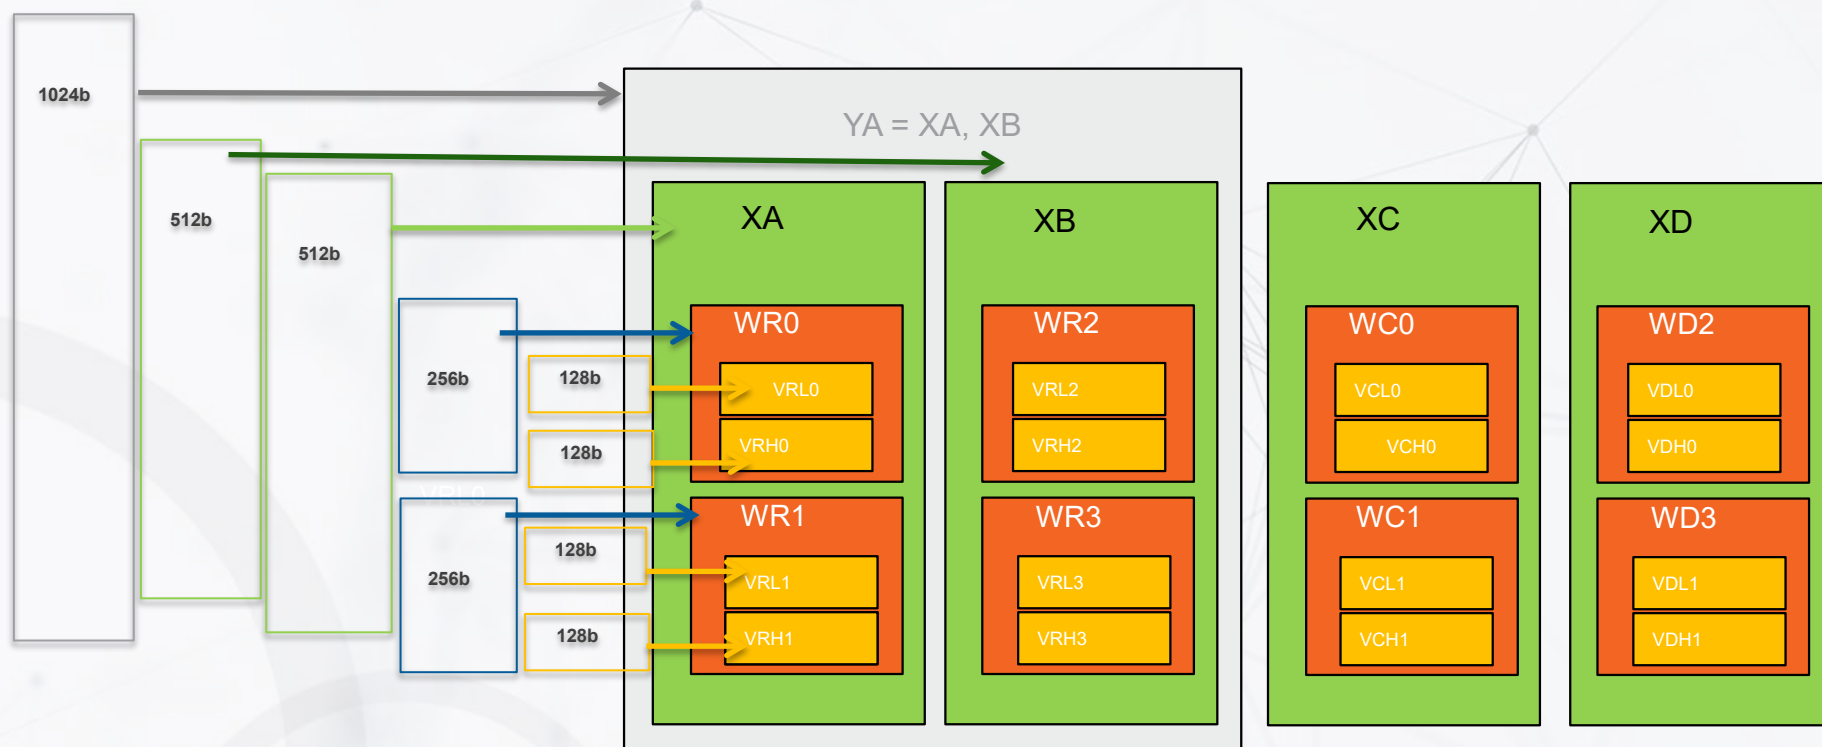
- Eight multiplications computed in parallel
- Vector elementary functions
  - Comparison, minimum, and maximum
- Fixed-point to floating-point conversion
- Floating-point to fixed-point conversion



# Vector Register File

High-width registers enabling SIMD instructions

- Incrementally grouped for higher widths

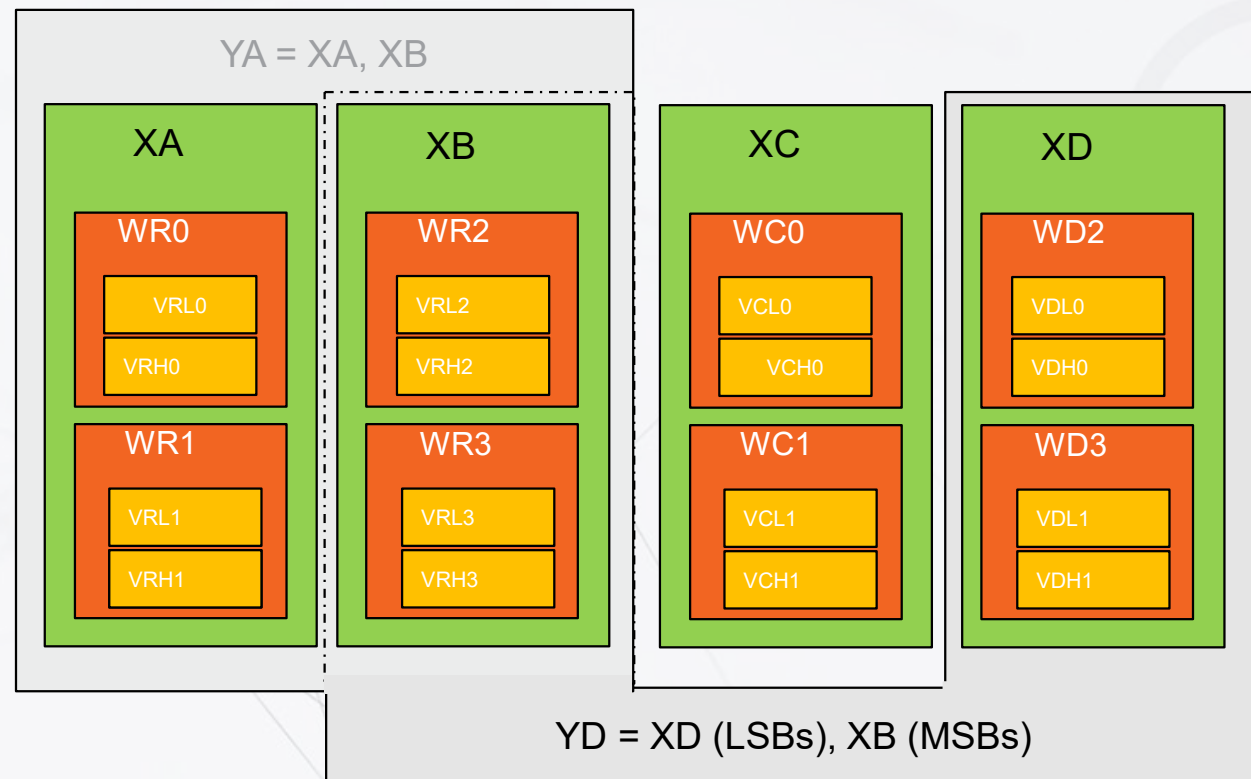




# Vector Register File


High-width registers enabling SIMD instructions

- Incrementally grouped for higher widths
- Basic hardware registers are 128-bits wide (V registers)



## Accumulator registers

- Store the results of the vector data path
- 8 vector lanes of 48 bits
- 32-bit mult. results and acc. without bit overflow



# Multi-precision Support

X Operand	Z Operand	Output	Number of GMACs @ 1 GHz
8 real	8 real	48 real	128
16 real	8 real	48 real	64
16 real	16 real	48 real	32
16 real	16 complex	48 complex	16
16 complex	16 real	48 complex	16
16 complex	16 complex	48 complex	8
16 real	32 real	48/80 real	16
16 real	32 complex	48/80 complex	8
16 complex	32 real	48/80 complex	8
16 complex	32 complex	48/80 complex	4
32 real	16 real	48/80 real	16
32 real	16 complex	48/80 complex	8
32 complex	16 real	48/80 complex	8
32 complex	16 complex	48/80 complex	4
32 real	32 real	80 real	8
32 real	32 complex	80 complex	4
32 complex	32 real	80 complex	4
32 complex	32 complex	80 complex	2
32 SPFP	32 SPFP	32 SPFP	8

## Example:

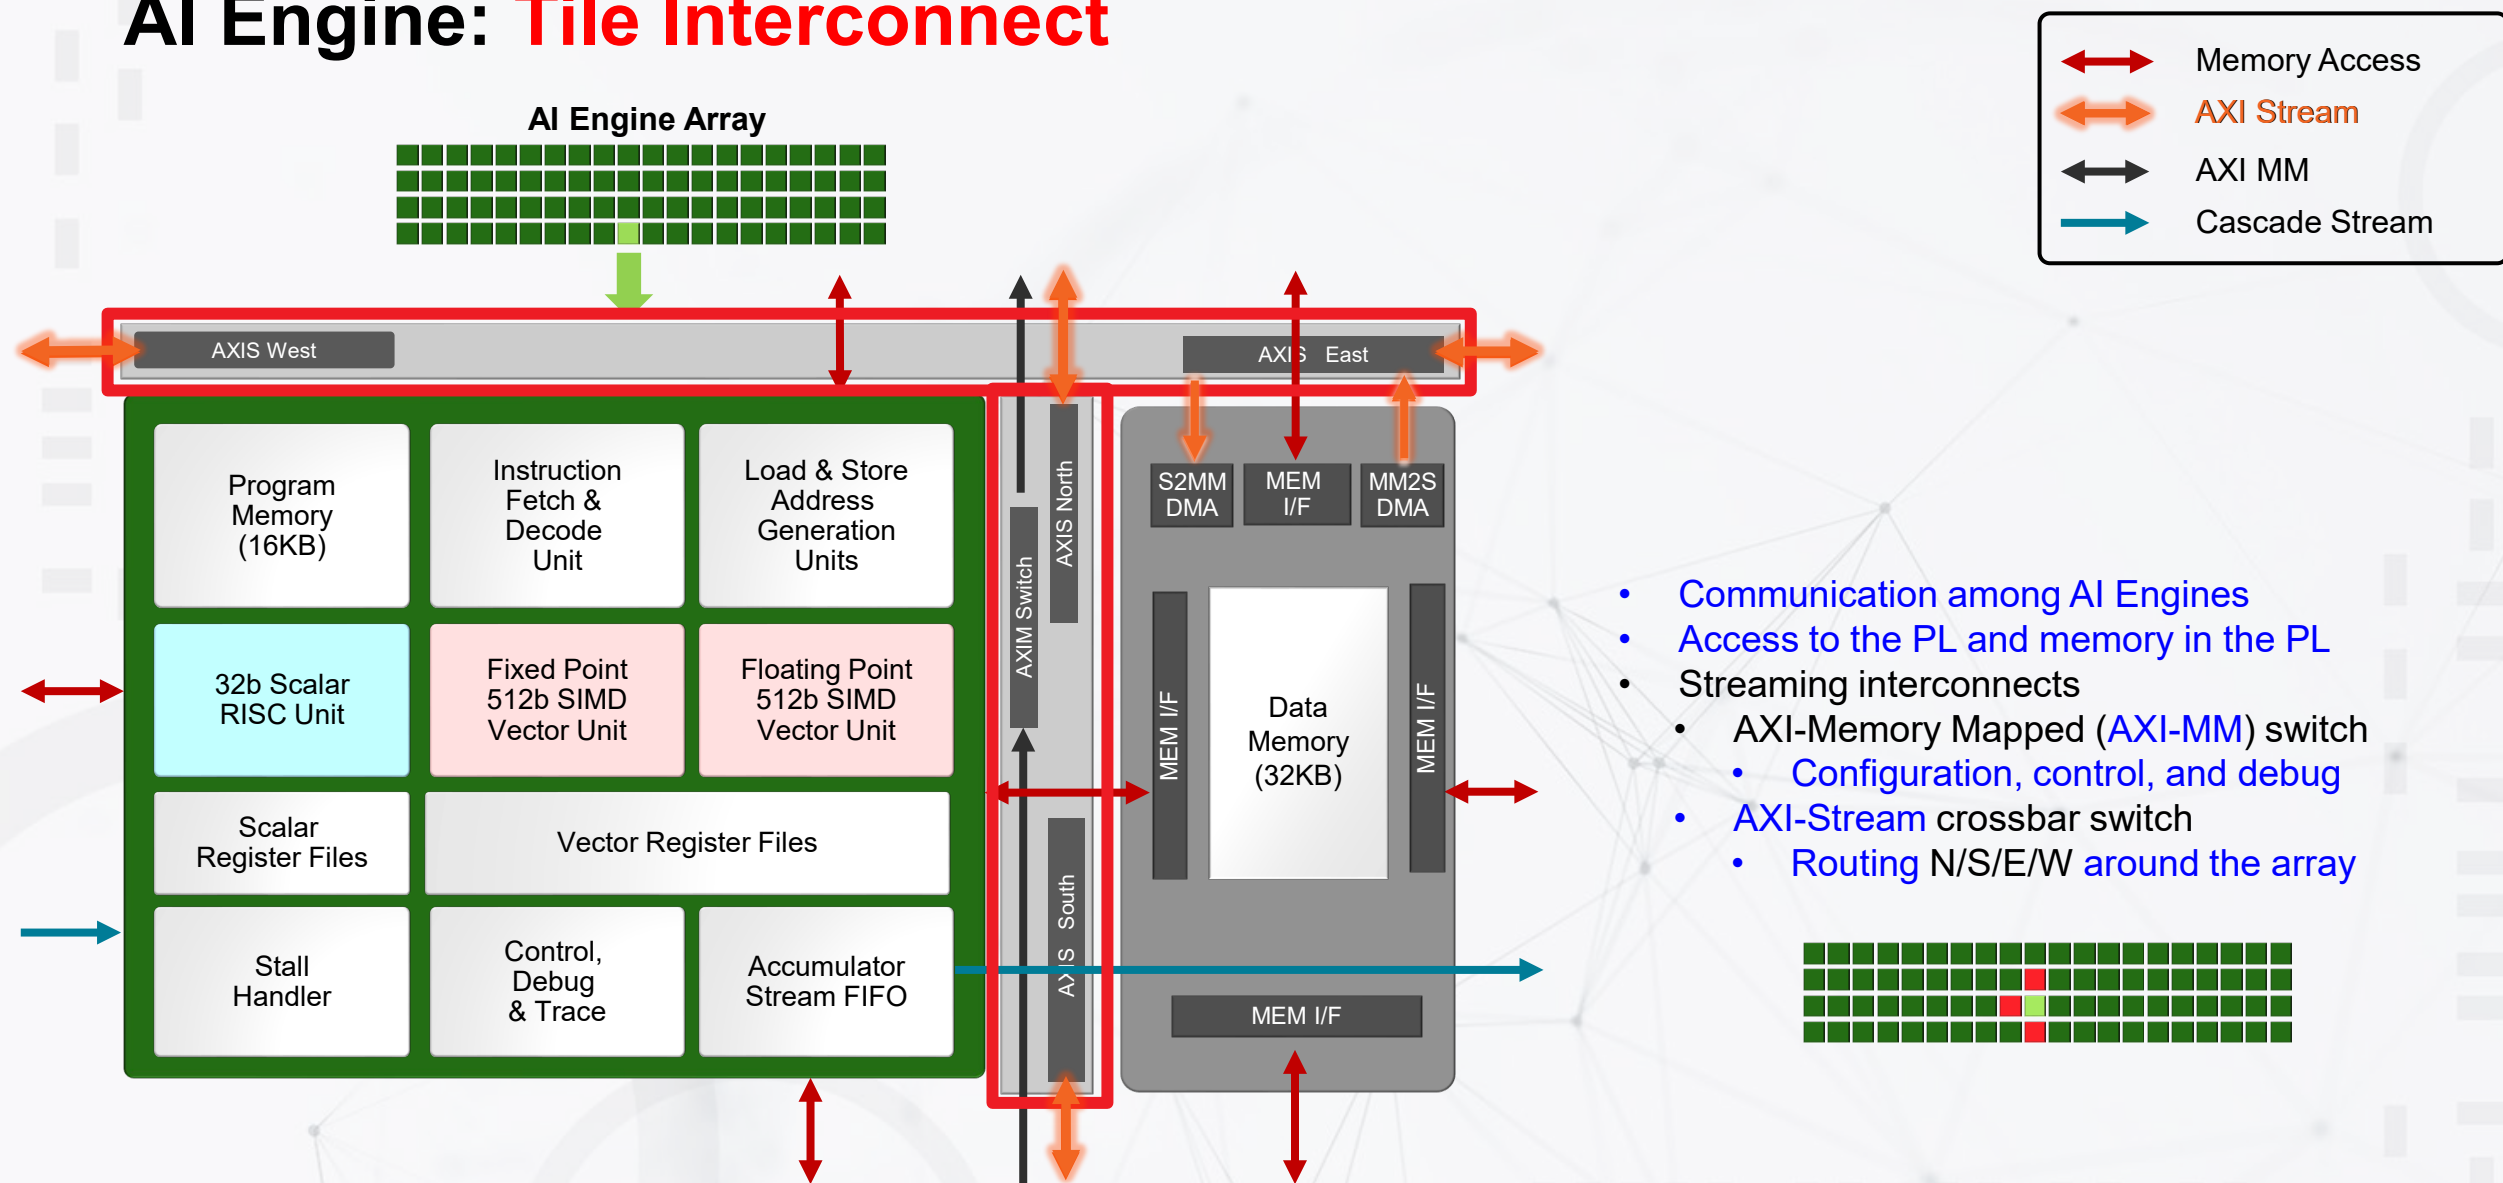
**X and Z => 16-bit input vectors**

No. of GMACs @ 1 GHz  
(32 MACs \* 1 GHz clock frequency) =  
**32 GMAC operations/second.**

**cfloat** is a vector type but is not directly supported by the AI Engine vector processor

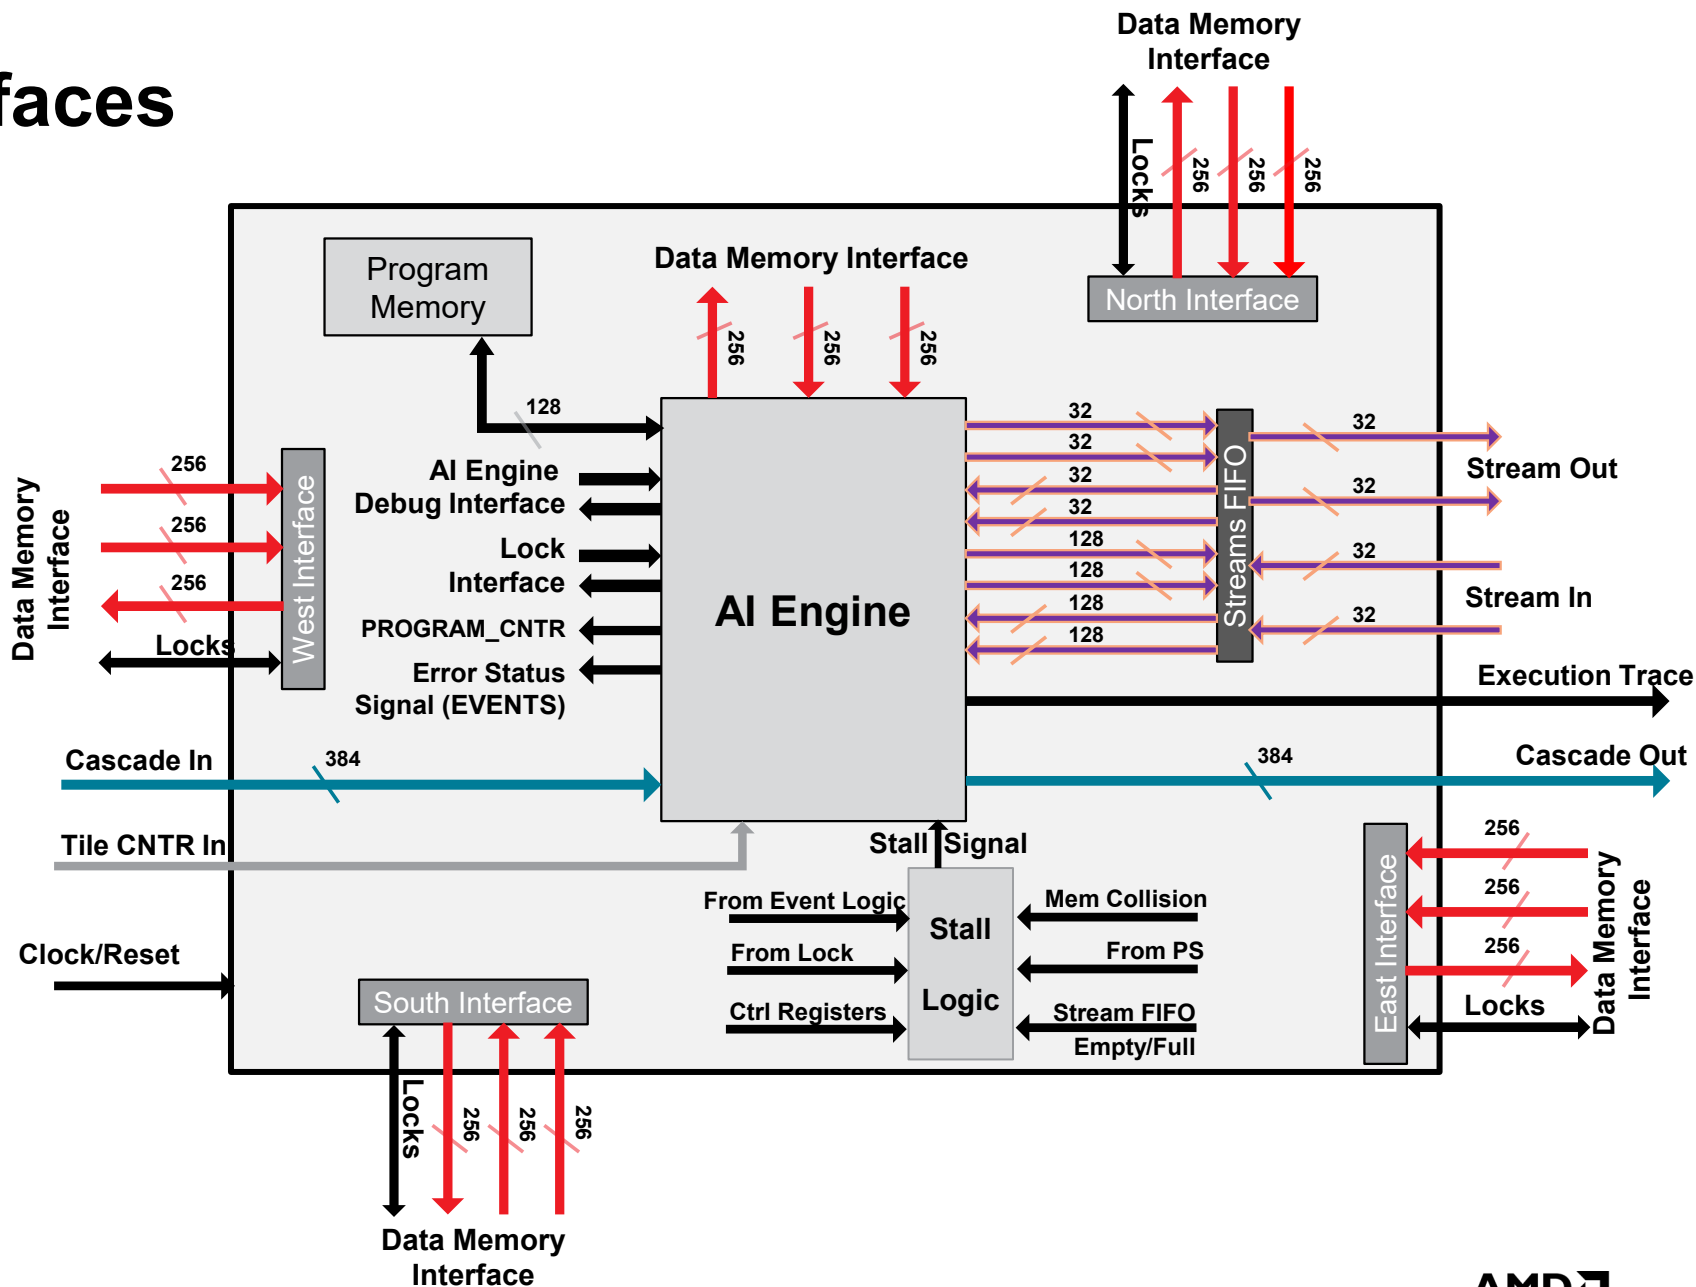
Two instructions have to be issued to perform a mult

# AI Engine: Tile Interconnect




# AI Engine Tile Interfaces

- Program Memory Interface
- Data Memory Interface
- Locks Interface
- Debug Interface
- Cascade Stream Interface
- Direct AXI4-Stream Interface
- Stall Handling Interface
- AI Engine Event Interface
- Tile Timer Interface
- Execution Trace Interface



## AI Engine: Memory





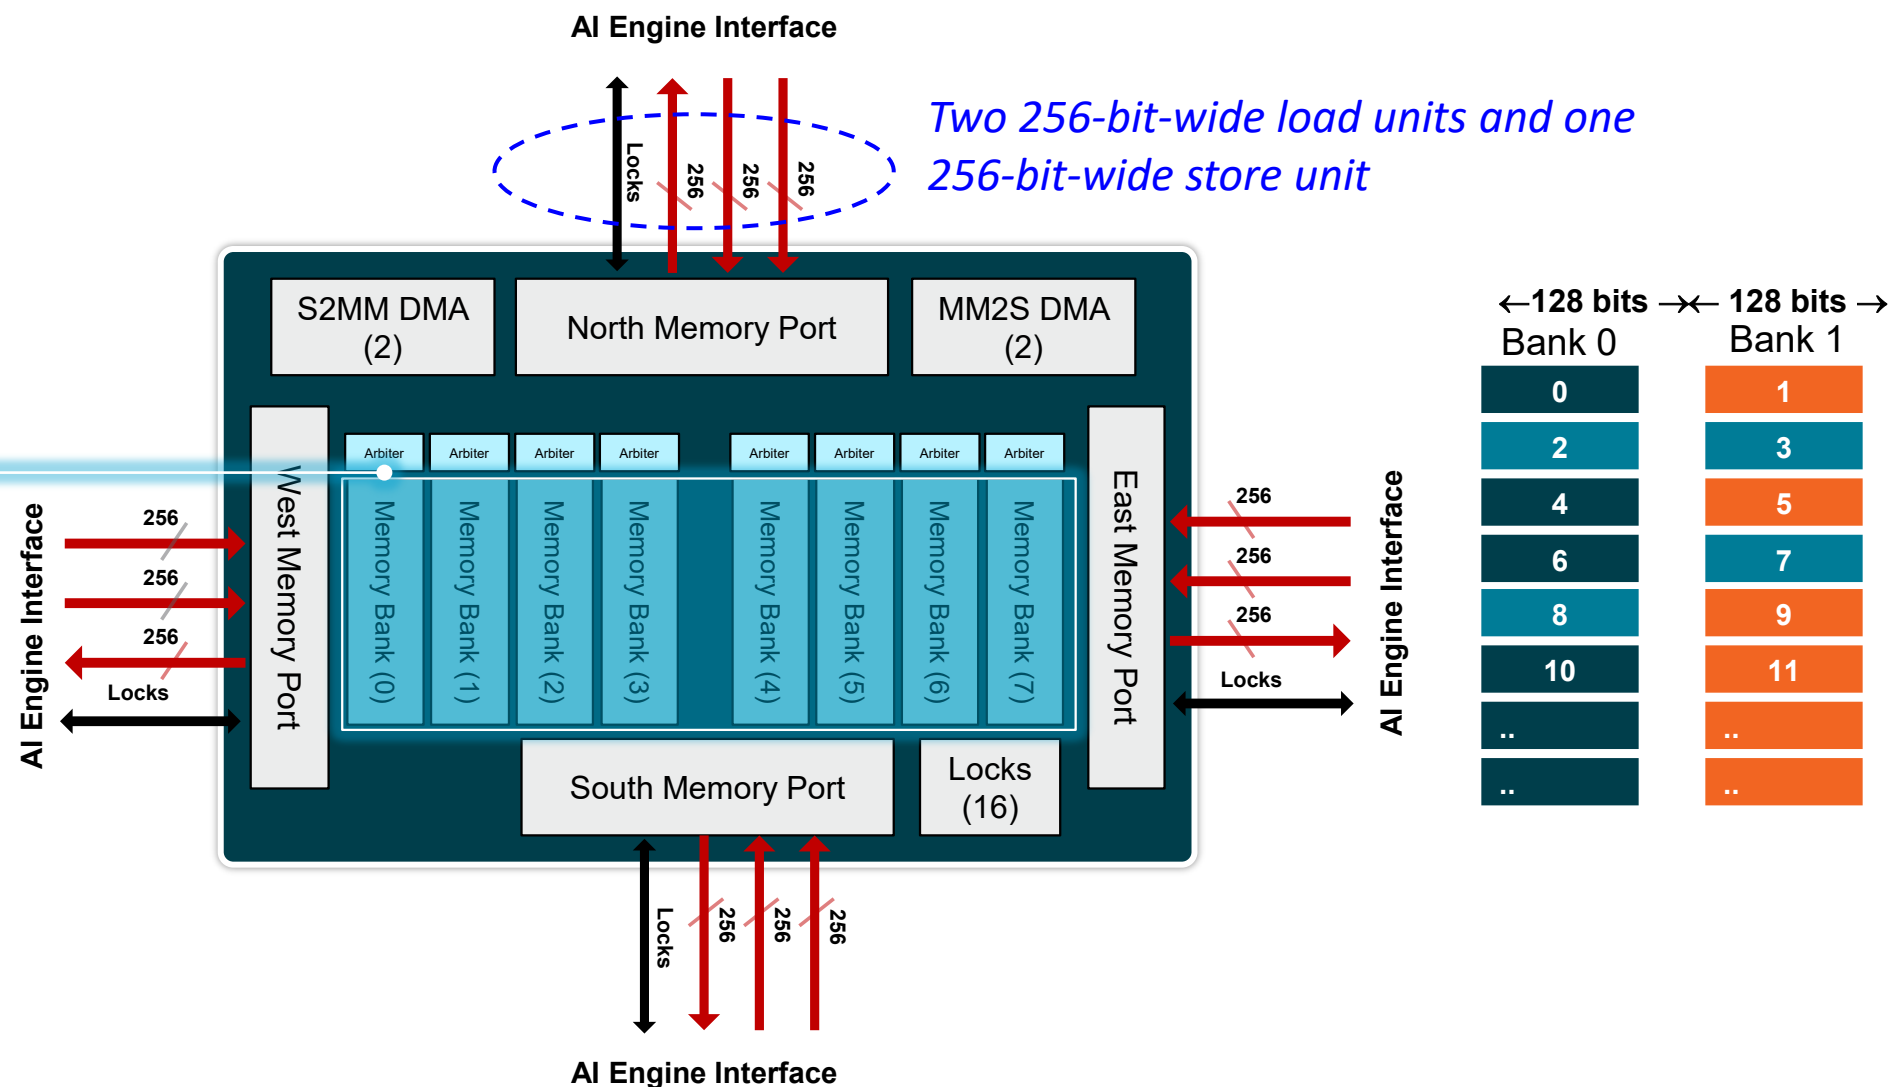
# AI Engine Data Memory: Banks

Eight memory banks:

- 256-word x 128-bit single-port memory
- Write enable for each 32-bit word
- Banks [0-1]: ECC protection
- Banks [2-7]: Parity check

ECC protection:


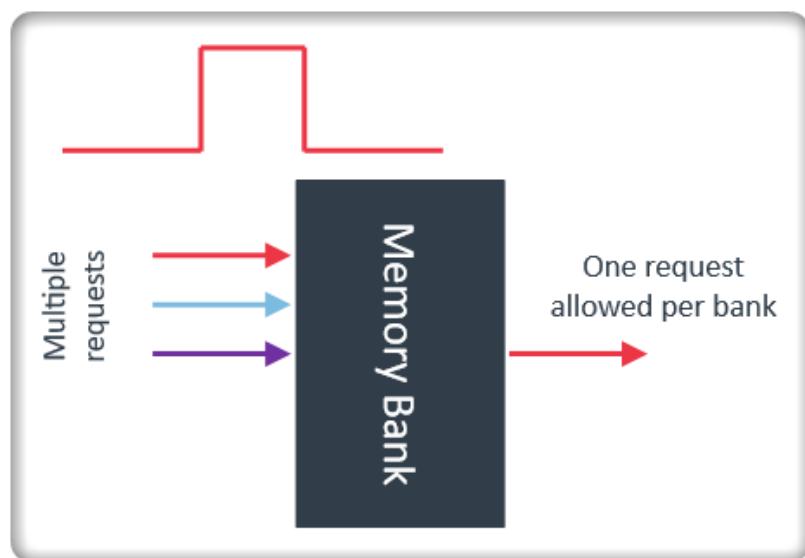
- 1-bit error detector/corrector
- 2-bit error detector per 32-bit word



# AI Engine Memory: Arbitration

## Memory Arbitration



- Round-robin to avoid starving any requester
- Handles a new request every cycle
- Only one requestor at a time allowed to access the memory per bank



# AI Engine Memory: Tile DMA Controller

## DMA Controller

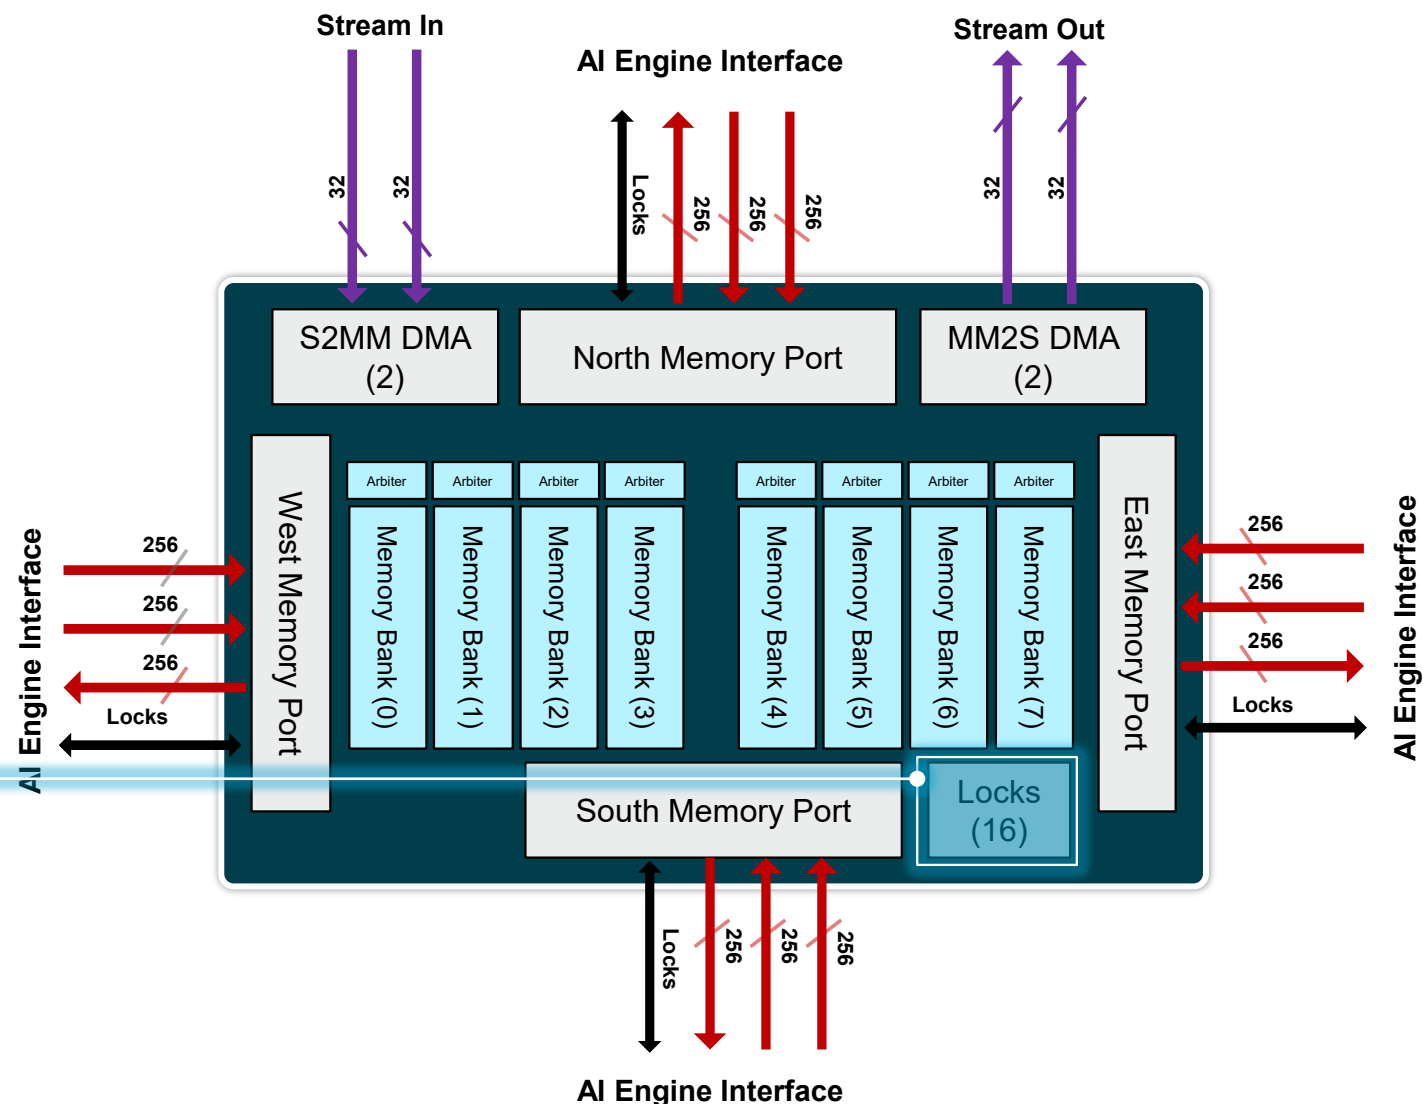
- Two incoming and two outgoing streams to the stream switches
- Divided into two modules
  - S2MM → store 32-bit streams of data to memory
  - MM2S → write the contents of the memory to a 32-bit stream
- Waits four clock cycles (128 bit) to access memory
- Has access to the 16 buffer descriptors
- Has access to the 16 locks



# AI Engine Memory: Lock Module

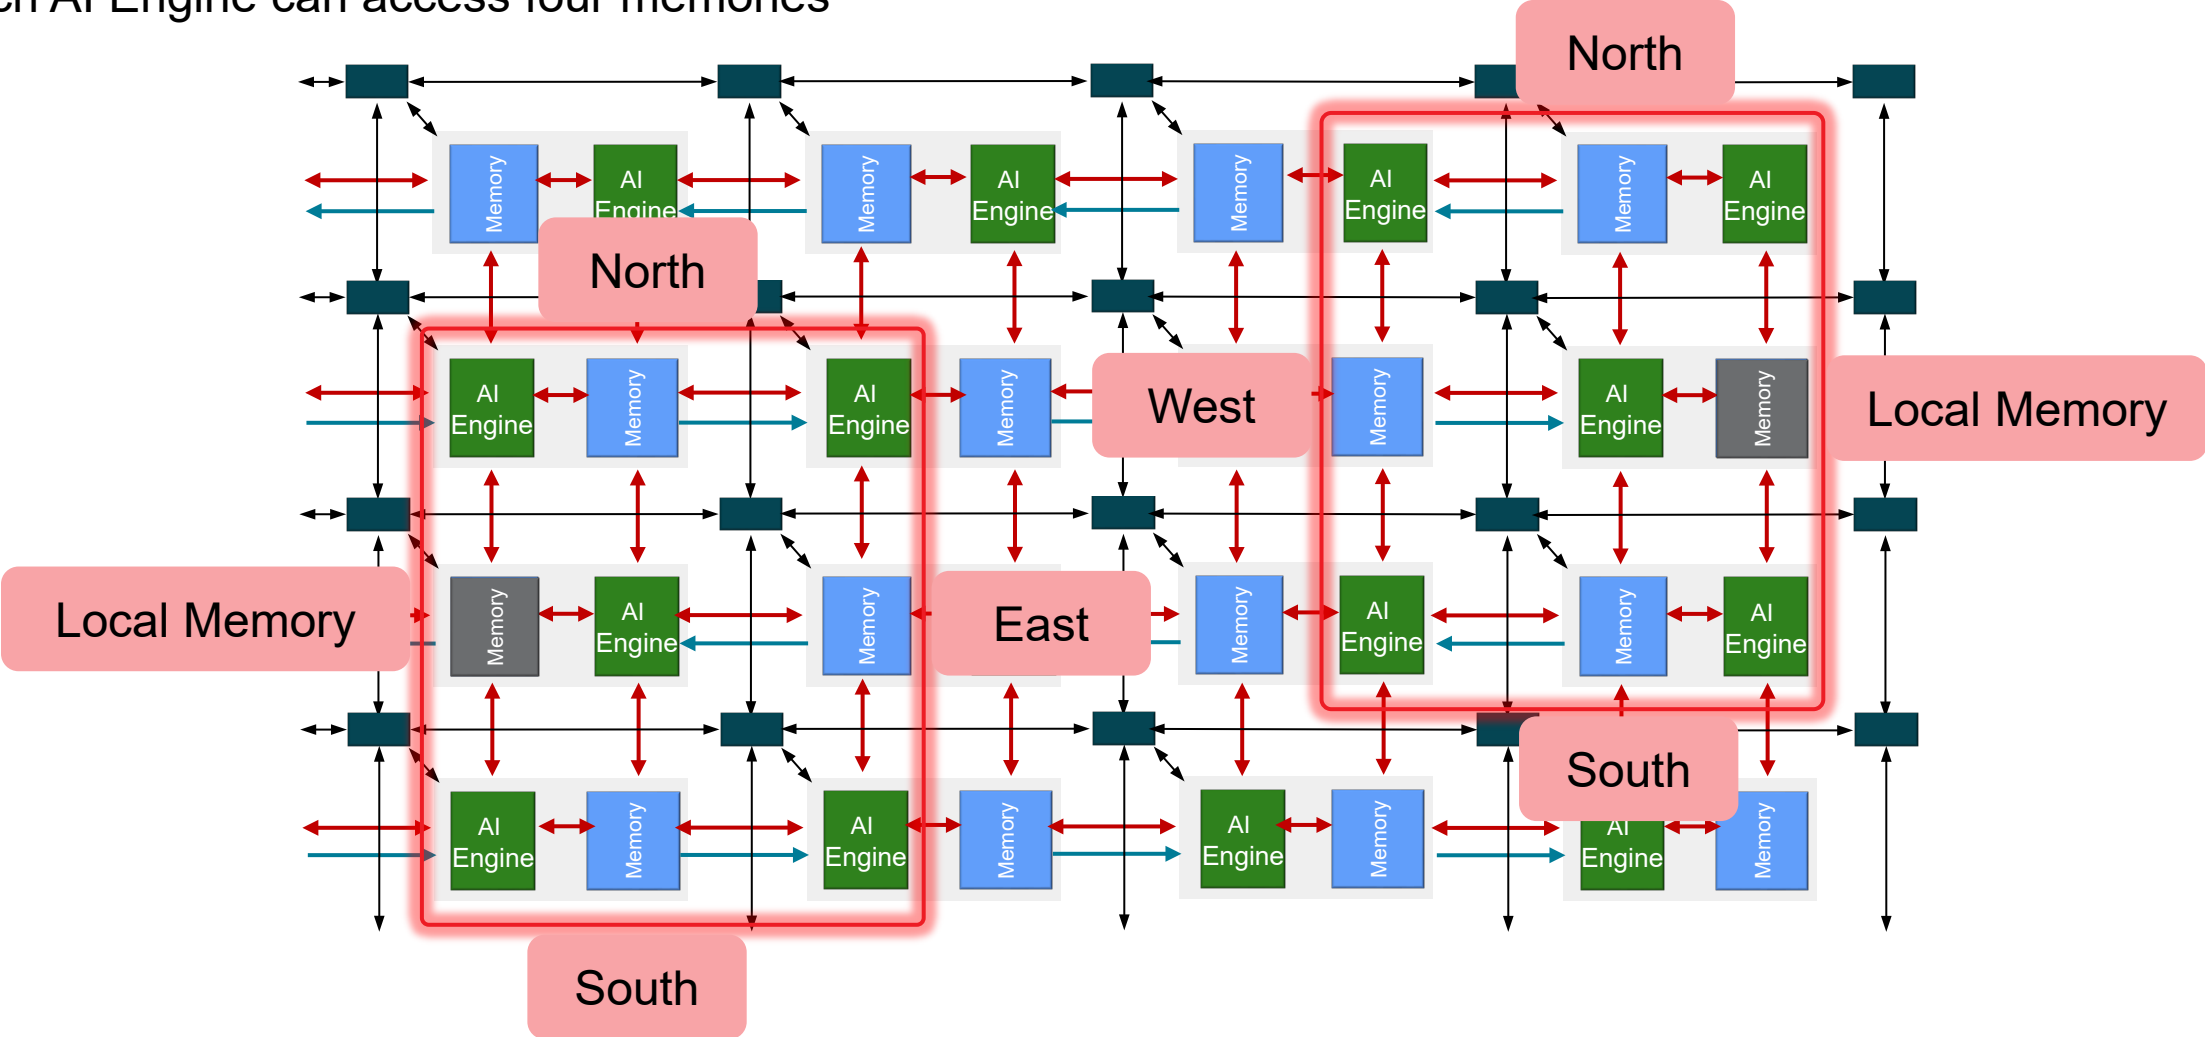
## Lock Module

- Helps achieve **synchronization among the AI Engines**, tile **DMA**, and external memory-mapped **AXI4 interface master**
- Sixteen hardware locks** with a **binary value** for each AI Engine memory module lock
- One arbitrator for each lock**
- Handles lock requests from the AI Engines, local DMA controller, and memory-mapped AXI4



# Memory Accesses in the AI Engine Array

Each AI Engine can access four memories




Result: All the diagonals do not have the same memory access structure

*circuit-switched / packet-switched*

# AXI4-Stream Interconnect

Each port handler contains 4-deep FIFO with 2-cycle latency

Each stream is capable of 4GB/s

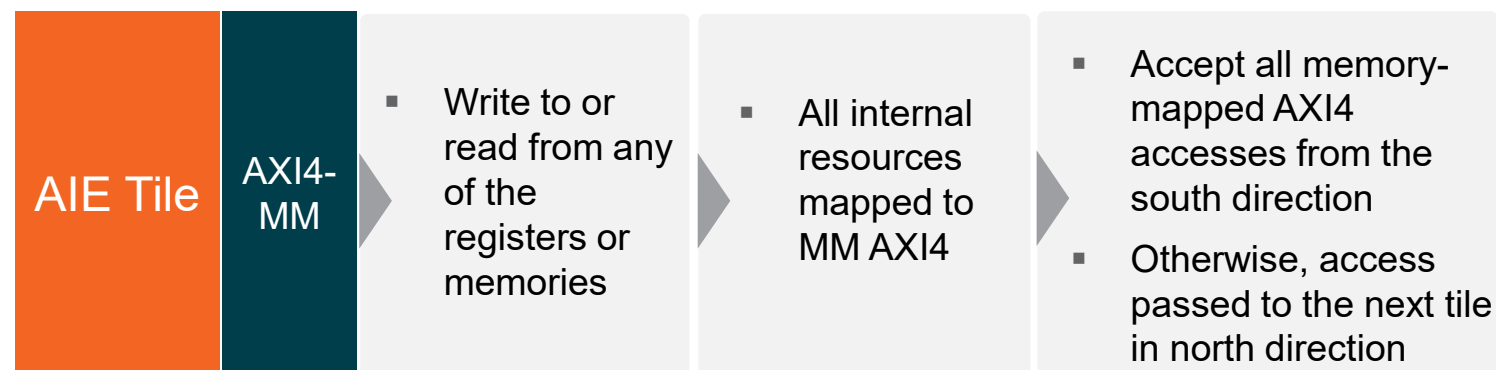


Interconnect crossing latency: 4 clock cycles

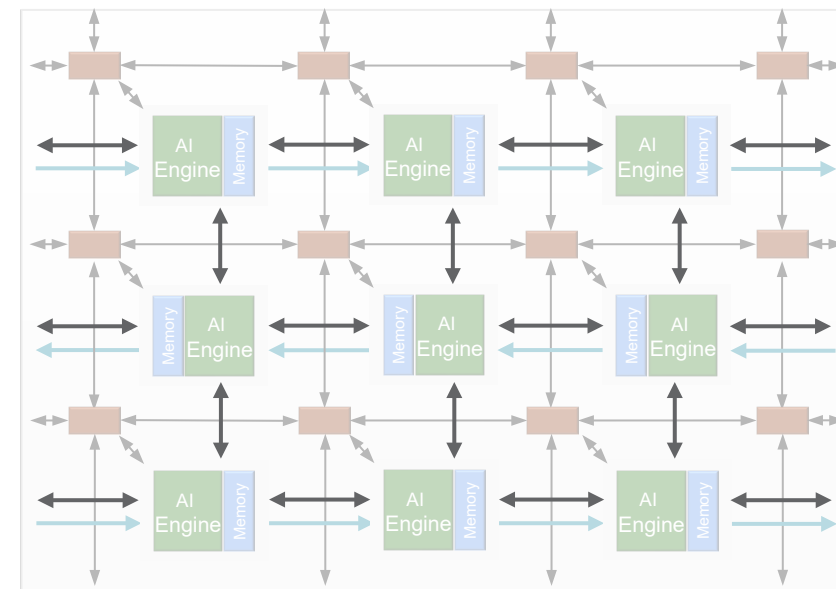


# AXI-MM Interconnect

The AI Engine MM AXI4 interconnect allows external access to internal AI Engine tile resources such as memories and registers for configuration and debug.

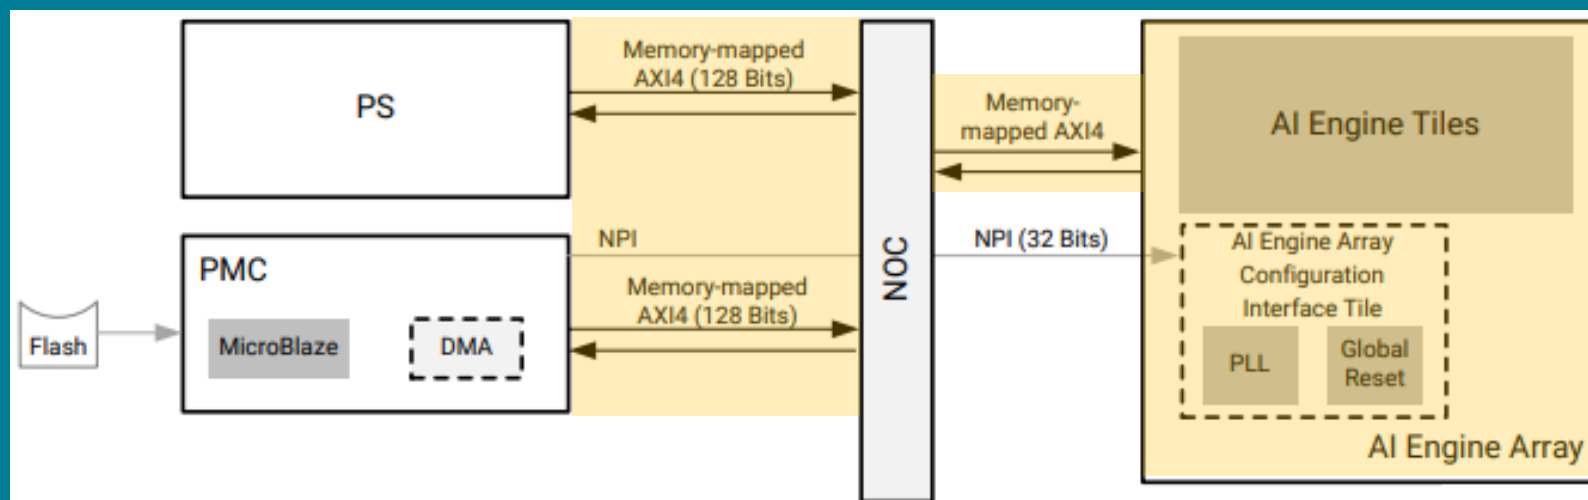


Memory Interface  
 Stream Interface  
 Cascade Interface

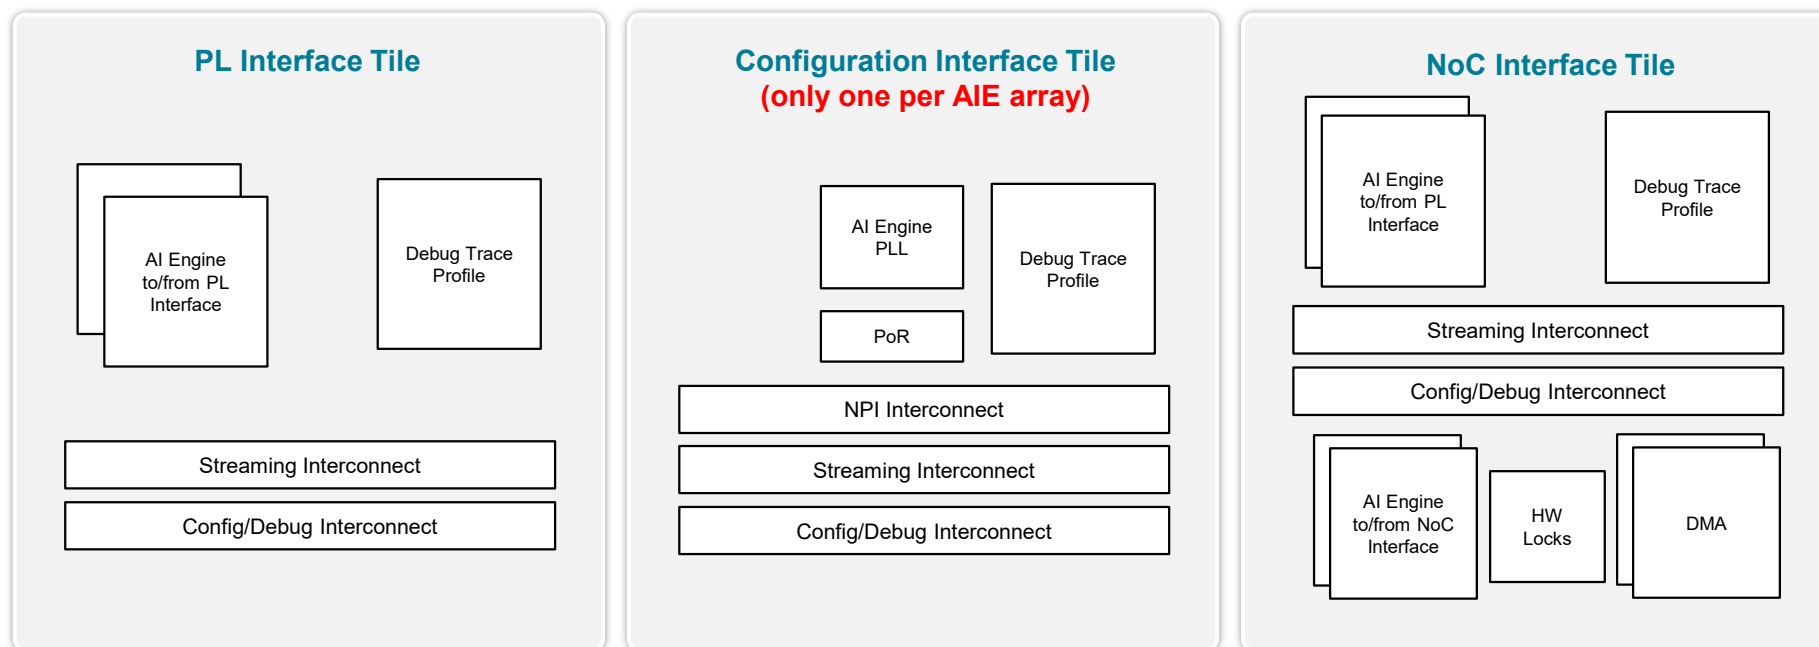


AI Engine Array

Driven from outside by any AXI4 master

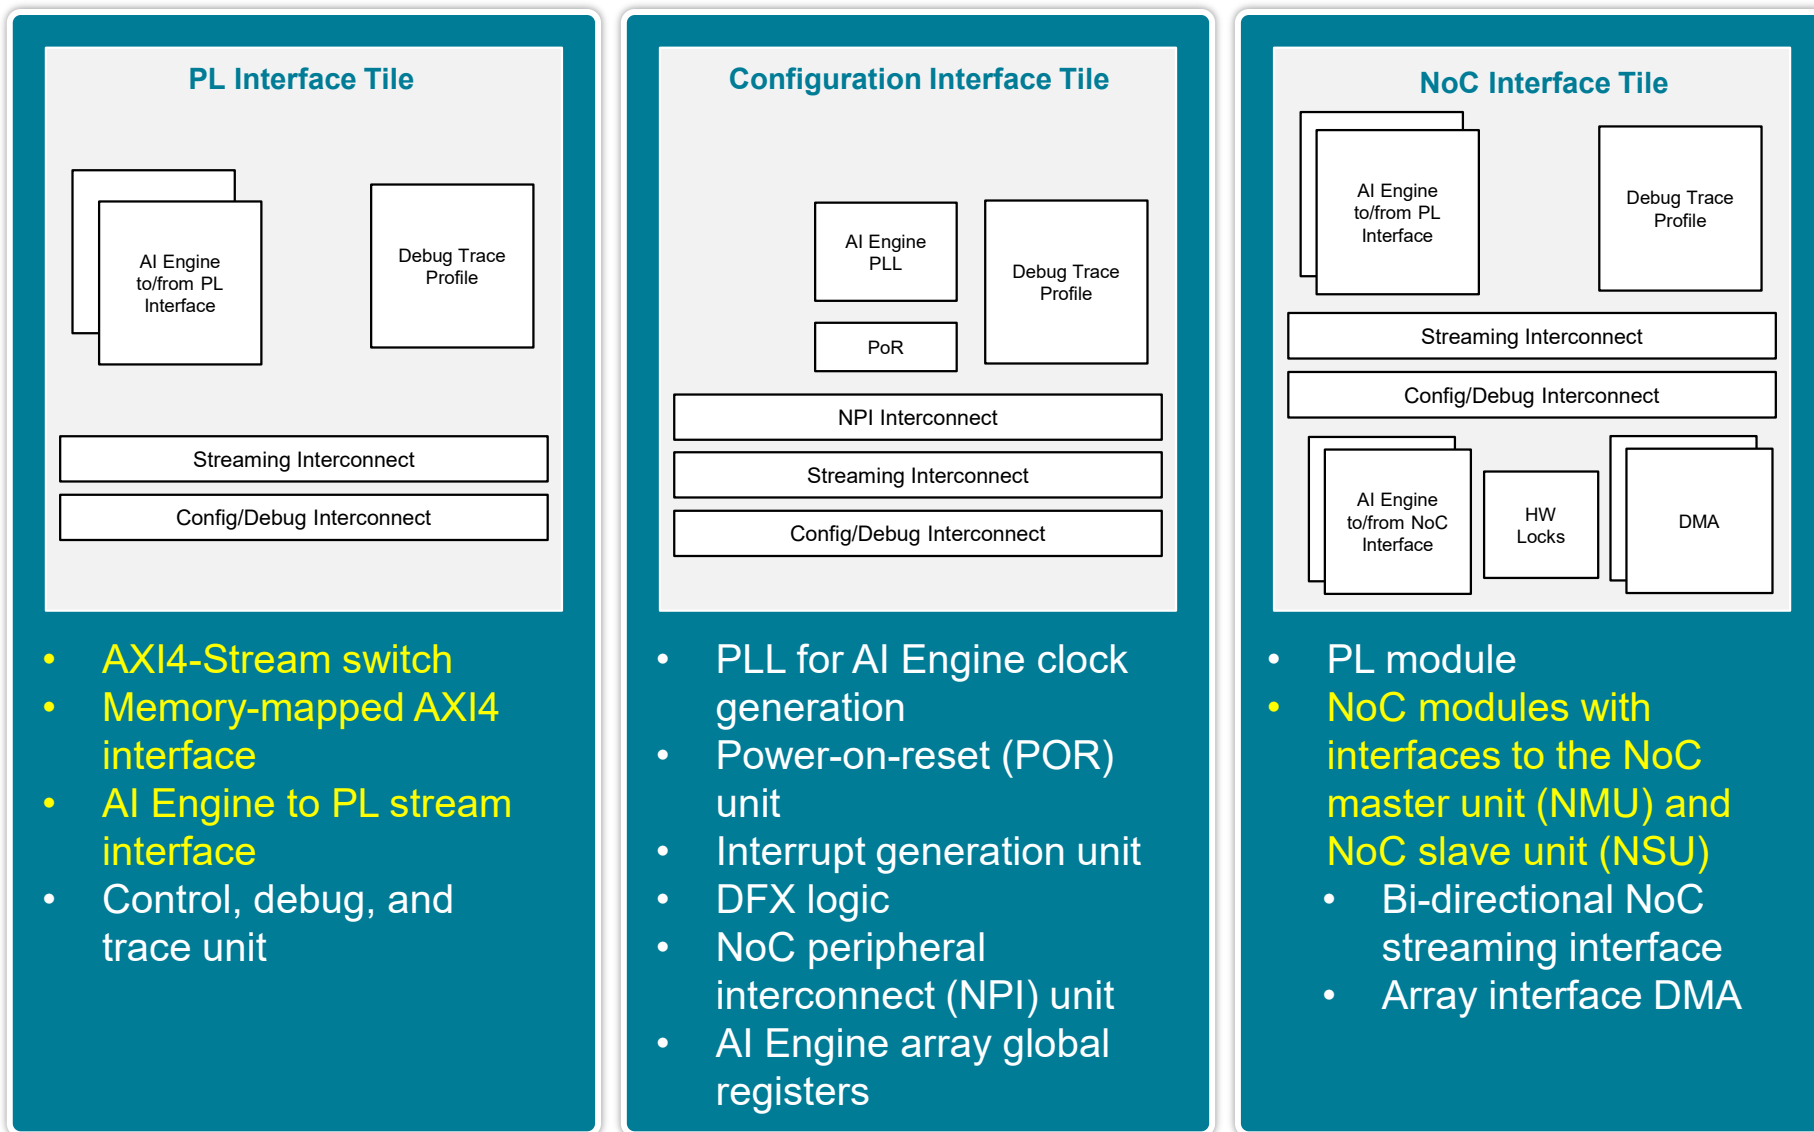


# AI Engine Array Interface Tiles (*shim*)



- One-to-one correspondence for every column of the AI Engine array
- Form a row and move memory-mapped AXI4 and AXI4-Stream data horizontally and vertically up
- Based on a modular architecture with device-specific final composition

# AI Engine Array Interface Tiles



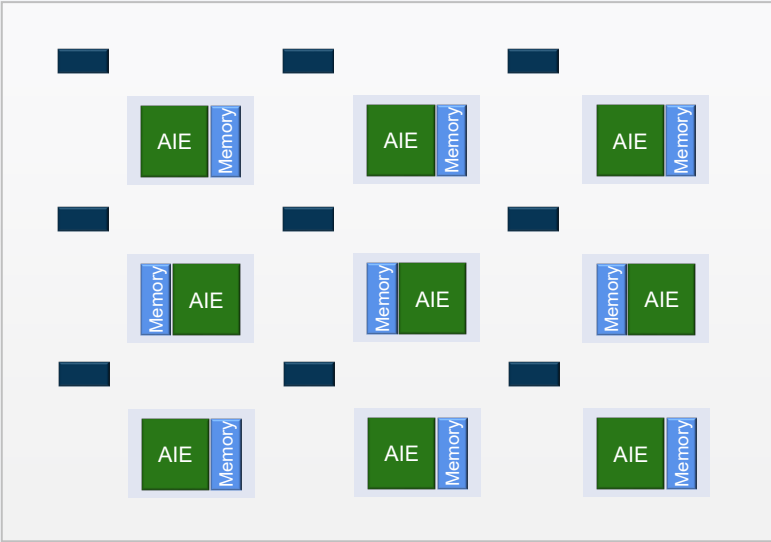
# Key Differences between AIE-ML and AIE

## AIE-ML features:

- Increased throughput for ML/AI inference workloads
- Optimized precision
- Increased on-chip memory
- Increased multiplier performance
- Power efficiency

	AIE	AIE-ML
Stream interface	Two 32-bit incoming and two 32-bit outgoing	One 32-bit incoming and one 32-bit outgoing
int32	Native support	No native support but can be emulated
FP32	Native support	No native support but can be emulated using bfloat16
BFLOAT16, INT4	Not supported	Supported
Memory tiles	N/A	Available

AIE 3x3 Array



AIE-ML 3x3Array



# Availability of AIE and AIE-ML

Versal™ series devices with the availability of AIE and AIE-ML

**Versal AI Core series:** (based on the devices, AIE or AIE-ML is available)

		VC1502	VC1702	VC1802	VC1902	VC2602	VC2802
AI Engines		198	304	300	400	0	0
AI Engines-ML		0	0	0	0	152	304
DSP Engines		1,032	1,312	1,600	1,968	984	1,312
AIE						AIE-ML	

**Versal AI Edge series:** (based on the devices, AIE or AIE-ML is available)

	VE2002	VE2102	VE2202	VE2302	VE1752	VE2602	VE2802
AI Engine-ML	8	12	24	34	0	152	304
AI Engines	0	0	0	0	304	0	0
DSP Engines	90	176	324	464	1,312	984	1,312
AIE-ML					AIE	AIE-ML	

**Versal Premium series:** (only AIE is available)

	VP1002	VP1052	VP1102	VP1202	VP1402	VP1502	VP2502	VP1552	VP1702	VP1802	VP2802	VP1902
AI Engines	-	-	-	-	-	-	472	-	-	-	472	-
DSP Engines	1,140	1,572	1,904	3,984	2,672	7,440	7,392	7,392	10,896	14,352	14,304	6,864
AIE							AIE					

# Summary

01

The hierarchy of the AI Engine architecture consists of:

- AI Engine array
- AI Engine tile
- AI Engine, memory module, and interconnects
- Scalar and vector units

02

The AI Engine tile has the following interfaces:

- Memory, lock, debug, stream, stall, event, timer, execution trace

03

The AI Engine memory module contains:

- Eight memory banks
- Two input streams to memory-map DMA (S2MM)
- Two memory-map to output DMA streams (MM2S)
- Hardware synchronization module (locks)



# GENERAL DISCLOSURE AND ATTRIBUTION STATEMENT

The information contained herein is for informational purposes only and is subject to change without notice. While every precaution has been taken in the preparation of this material, it may contain technical inaccuracies, omissions and typographical errors, and AMD is under no obligation to update or otherwise correct this information. Advanced Micro Devices, Inc. makes no representations or warranties with respect to the accuracy or completeness of the contents of this material, and assumes no liability of any kind, including the implied warranties of noninfringement, merchantability or fitness for particular purposes, with respect to the operation or use of AMD hardware, software or other products described herein. No license, including implied or arising by estoppel, to any intellectual property rights is granted by this document. Terms and limitations applicable to the purchase or use of AMD's products are as set forth in a signed agreement between the parties or in AMD's Standard Terms and Conditions of Sale. GD-18.

©2024 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, MicroBlaze, Versal, and combinations thereof are trademarks of Advanced Micro Devices, Inc. Arm and Cortex are registered trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. PCIe is a registered trademark of PCI-SIG Corporation. Other product names used in this publication are for identification purposes only and may be trademarks of their respective owners.

